

COMPUTER SCIENCE



Database Management System

ER Model



Lecture_1

Vijay Agarwal sir

TOPICS
TO BE
COVERED

01

Time Stamp Protocol

02

ER MODEL






Time Stamp Protocol . ① T_i : issue Read(Q)

$TS(T_i)$

Data Item $\rightarrow RTS(Q)$
Time Stamp $\rightarrow WTS(Q)$

$TS(T_i) < WTS(Q)$; Reject





TIMESTAMP BASED CONCURRENCY CONTROL

Timestamp-Based Protocols

- ❑ Each transaction T_i is issued a timestamp $TS(T_i)$ when it enters the system.
 - ❖ Each transaction has a unique timestamp
 - ❖ Newer transaction have timestamp strictly greater than earlier ones
 - ❖ Timestamp could be based on a logical counter
 - Real time may not be unique
 - Can use (wall-clock time, logical counter) to ensure
- ❑ Timestamp-based protocols manage concurrent execution such that **time-stamp order = serializability order**
- ❑ Several alternative protocols based on timestamps

Timestamp-Based Protocols

The **timestamp ordering (TSQ) protocol**

- ❑ Maintains for each data Q two timestamp values:
 - ❖ W-timestamp(Q) is the largest time-stamp of any transaction that executed write (Q) successfully.
 - ❖ R-timestamp (Q) is the largest time-stamp of any transaction that executed **read** (Q) successfully.

Note

- ❑ Imposes rules on read and write operations to ensure that
 - ❖ any conflicting operations are executed in timestamp order
 - ❖ out of order operations cause transaction rollback

$T_i : \text{Read}(Q)$

- ① $TS(T_i) < WTS(Q)$ Reject & Rollback

$T_i : \text{Write}(Q)$

- ① $TS(T_i) < RTS(Q)$
② $TS(T_i) < WTS(Q)$ } Write operation
Reject & Rollback

Conflict operation

$R^{T_i}(Q) - W^{T_j}(Q)$

$W^{T_i}(Q) - R^{T_j}(Q)$

$W^{T_i}(Q) - W^{T_j}(Q)$

V.V.V.V. Imp.

Note

All conflicting operation
order must be same as

Transaction order (Timestamp
order)



I: T_i - Read(Q) (Transaction T_i Issue R(Q) Operation)

- (i) If $TS(T_i) < WTS(Q)$: Read operation Reject & T_i Rollback.
- (ii) If $TS(T_i) \geq WTS(Q)$: Read operation is allowed
and Set Read - $TS(Q) = \max[RTS(Q), TS(T_i)]$

CHECK
 $R(Q) \vdash WTS(Q)$
 $W(Q) \vdash RTS(Q)$
 $W(Q) \vdash WTS(Q)$

II: T_i - Write(Q) (Transaction T_i Issue Write(Q) Operation)

- (i) If $TS(T_i) < RTS(Q)$: Write operation Reject & T_i Rollback.
- (ii) If $TS(T_i) < WTS(Q)$: Write operation Reject & T_i Rollback.
- (iii) Otherwise execute write (Q) operation
Set Read $WTS(Q) = TS(T_i)$

2020	2021	2023
T_1	T_2	T_3
	$R(Q)$	$\underline{w(Q)}$
		mill

$$TS(T_2) < wTS(Q)$$

$$\underline{20 < 30}$$

Timestamp-Based Protocols (Cont.)

- ❑ Suppose a transaction T_i issues a **read** (Q)
 1. If $TS(T_i) \leq \text{W-timestamp}(Q)$, then T_i needs to read a value of Q that was already overwritten.
 - Hence, the **read** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) \geq \text{W-timestamp}(Q)$, then the **read** operation is executed, and $R\text{-timestamp}(Q)$ is set to.

$$\max(R\text{-timestamp}(Q), TS(T_i)).$$

Timestamp-Based Protocols (Cont.)

- Suppose a transaction T_i issues **write**(Q)
 1. If $TS(T_i) < \text{R-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that the value would never be produced.
 - Hence, the **write** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) < \text{W-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q.
 - Hence, this **write** operation is rejected, and T_i is rolled back.
 3. Otherwise, the **write** operation is executed, and $\text{W-timestamp}(Q)$ is set to $TS(T_i)$.

10 T_1	T_i 20 T_2	30 T_3
	<u>$W(Q)$</u>	
		<u>$R(Q)$</u>

	20 T_2	30 T_3
	<u>$W(Q)$</u>	
		<u>$W(Q)$</u>

I

① $T_1 \Rightarrow R(A)$

$W.TS(A) = 0$

$TS(T_1) = 10$

T_1 Read allowed.

② $T_2 \Rightarrow W(A)$
(20)

$R(A) = 10$

$W(A) = 0$

T_2 Write allowed

$\begin{cases} TS(T_2) < RTS(A) \\ TS(T_2) < WTS(A) \end{cases}$

③ ⑩ $T_1 \Rightarrow W(A)$ Issue $RTS(A) = 10$
 $WTS(A) = 20$
 $TS(T_1) < WTS(A)$
 $10 < 20$ Reject, Not TSP

$TS(T_1) = 10$

$TS(T_2) = 20$


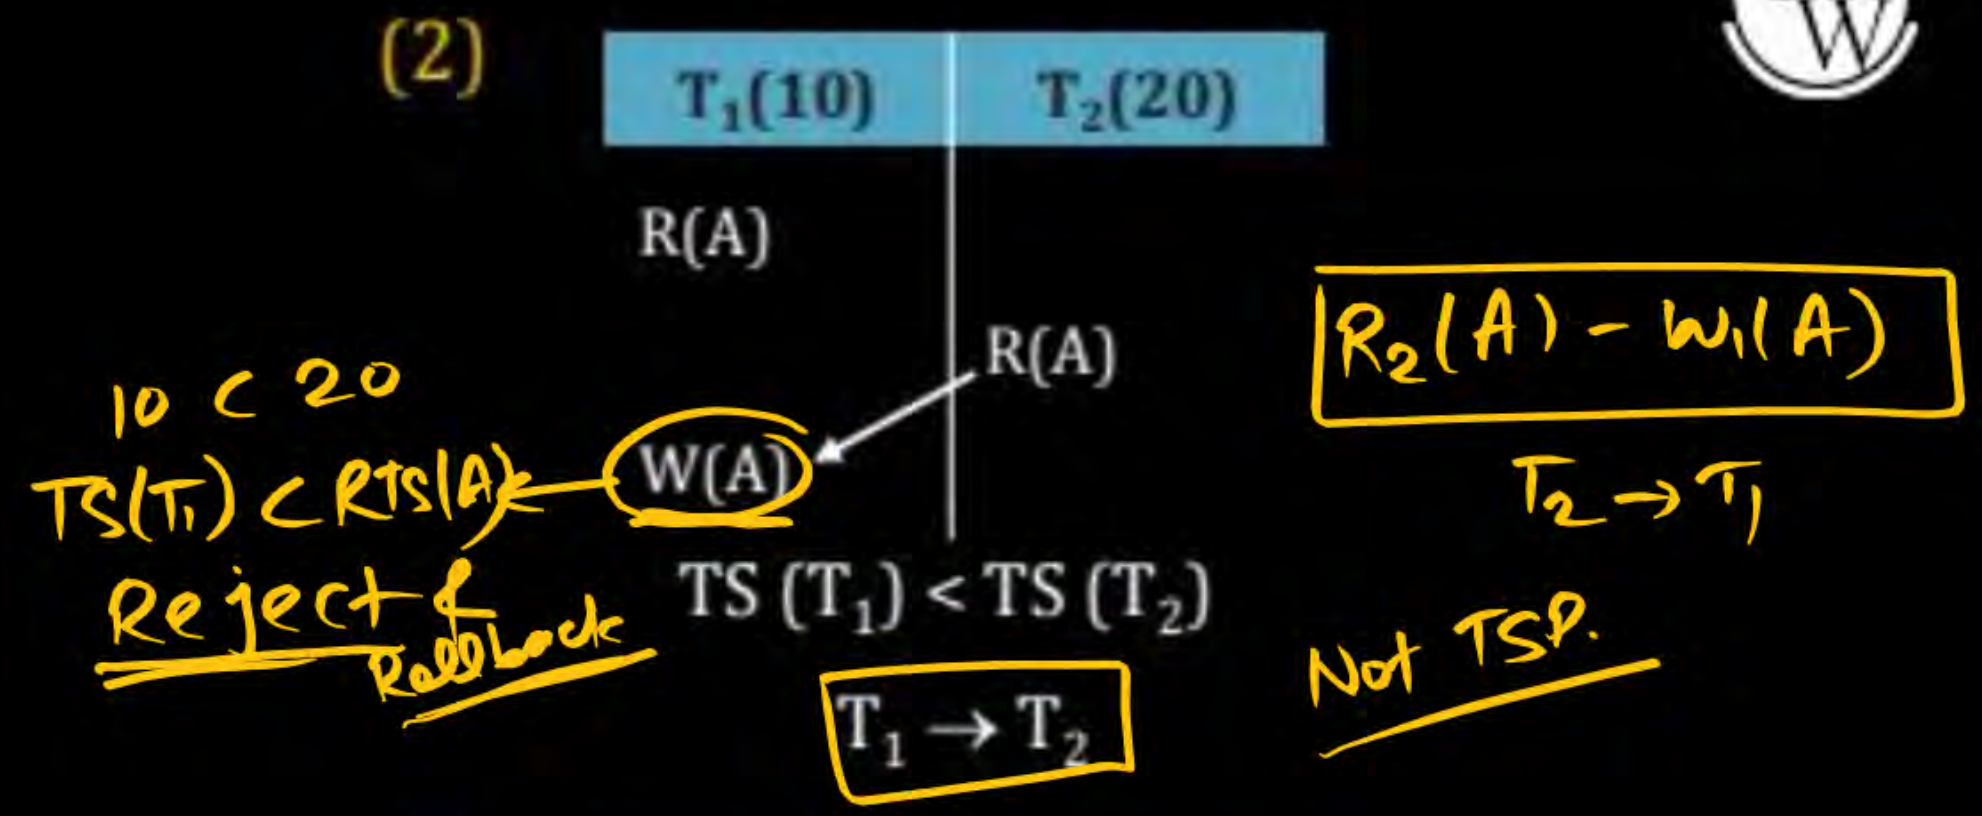

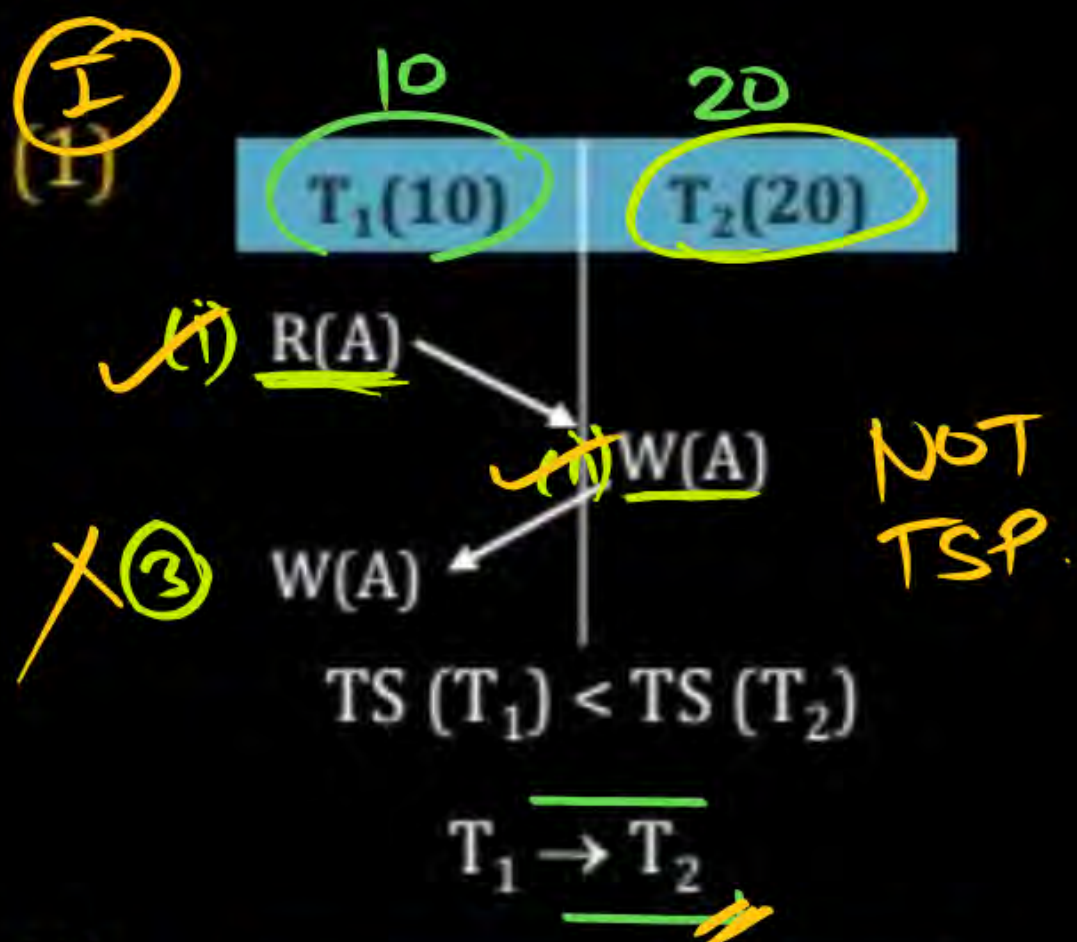
$TS(T_1) < TS(T_2)$

$T_1 \rightarrow T_2$

All conflict operation
order must be $T_1 \rightarrow T_2$.

$R(A) - W(A) \Rightarrow T_1 \rightarrow T_2$

$W_2(A) - W_1(A) = T_2 \rightarrow T_1$ X Not TSP.



Thomas Write Rule (View Serializable)

$T_i: \text{Read}(Q)$

$TS(T_i) < WTS(Q)$; Reject & Rollback Same as TSP

$T_i: \text{Write}(Q)$

$TS(T_i) < RTS(Q)$; Reject & T_i Rollback. Same as TSP

$TS(T_i) < WTS(Q)$; Ignore, No Rollback

Thomas' Write Rule



- ❑ Modified version of the timestamp-ordering protocol in which **obsolete write** operations may be ignored under certain circumstances.
- ❑ When T_i attempts to write data item Q , if $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of $\{Q\}$.
 - ❖ Rather than rolling back T_i as the timestamp ordering protocol would have don, this **{write}** operation can be ignored.
- ❑ Otherwise this protocol is the same as the timestamp ordering protocol.
- ❑ Thomas' Write Rule allows greater potential concurrency.
 - ❖ Allows some view-serializable schedules that are not conflict-serializable.

Thomas Write Rule (View Serializability)



1. $TS(T_i) < RTS(Q)$: Rollback *Reject*
2. $TS(T_i) < WTS(Q)$ Write operation is Ignored and No Roll back

Same as TSP

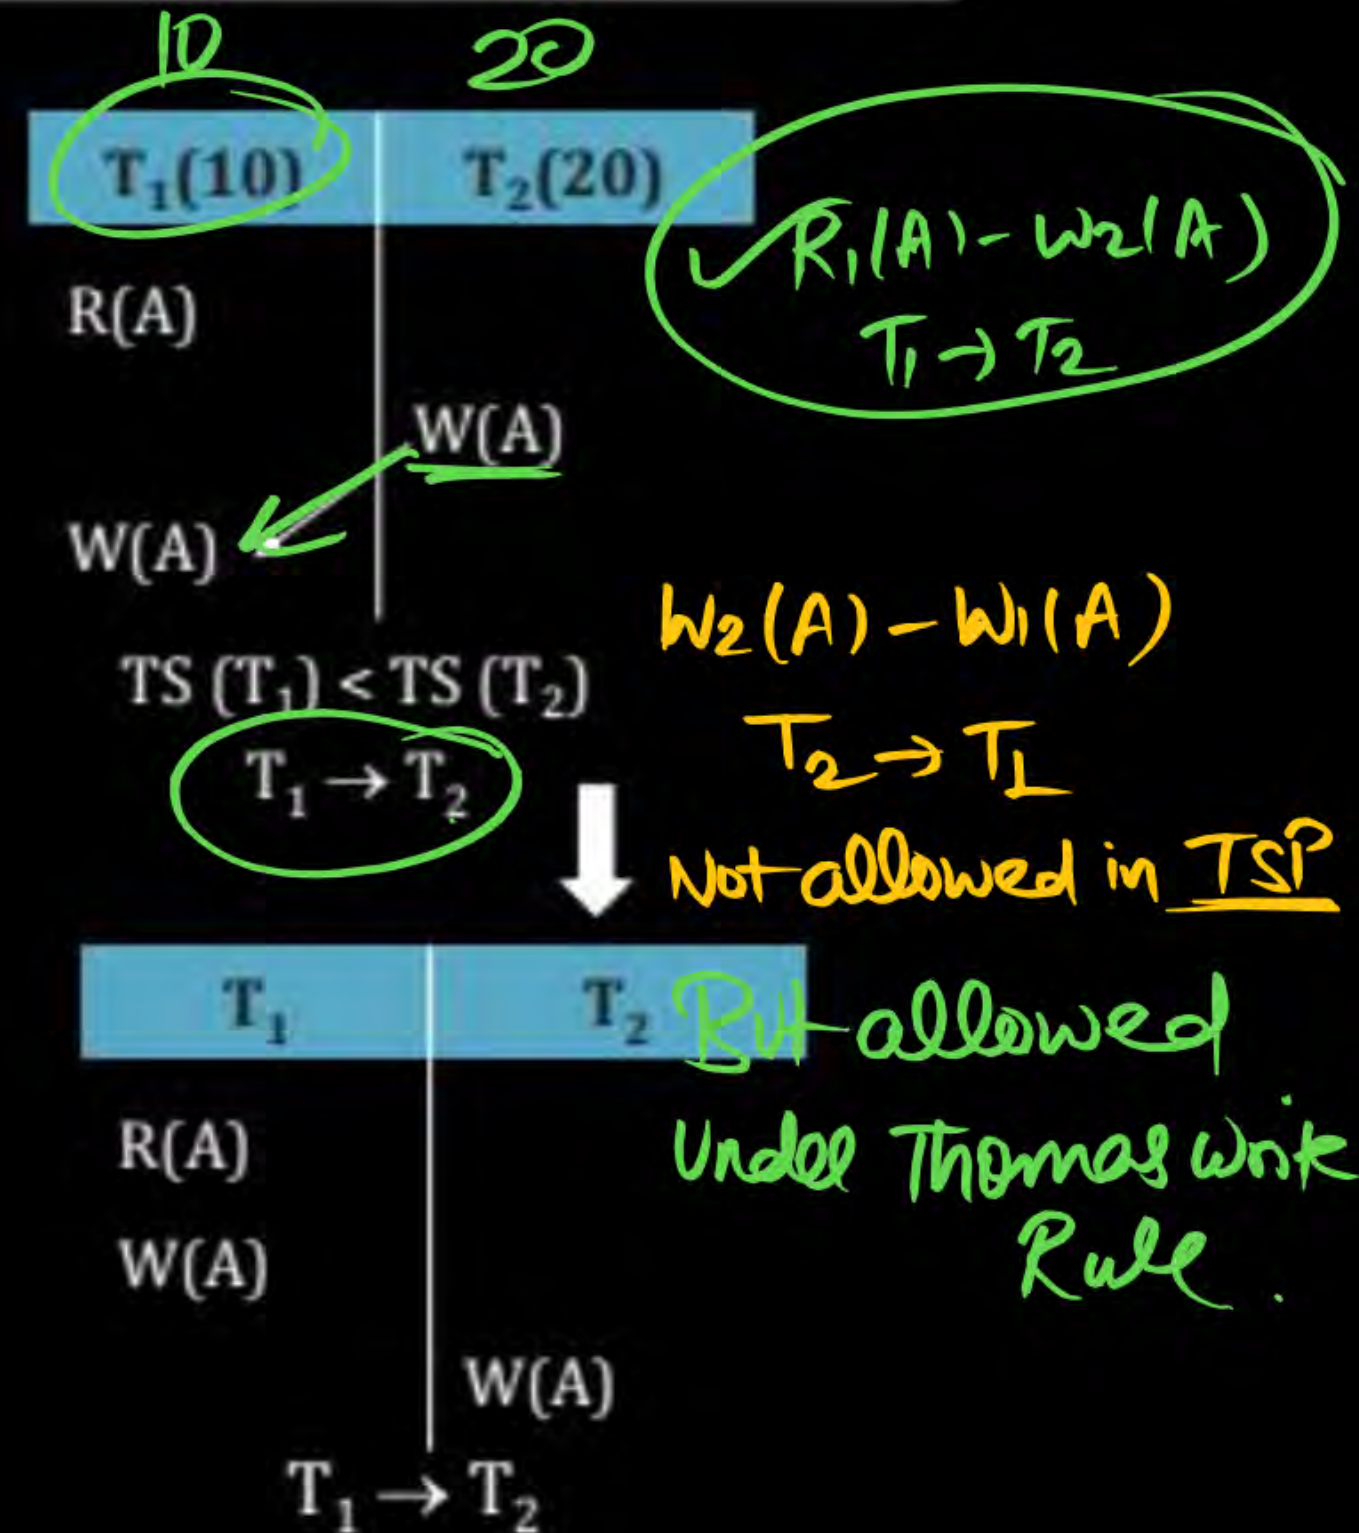
Time Stamp Protocol: Ensure serializability deadlock free but starvation possible

Deadlock Prevention Algorithm

(1) Wait-Die

(2) Wound-wait

Older Younger



Q.1



Consider the following database schedule with two transactions, T_1 and T_2 .

$S = r_2(X); r_1(X); r_2(Y); w_1(X); r_1(Y); w_2(X); a_1; a_2$

where $r_i(Z)$ denotes a read operation by transaction T_i on a variable Z , $w_i(Z)$ denotes a write operation by T_i on a variable Z and a_i denotes an abort by transaction T_i

Which one of the following statements about the above schedule is TRUE?

[MCQ:2016-2M]

- ☒ A S is non-recoverable
- ☒ B S is recoverable, but has a cascading abort
- ☒ C S does not have a cascading abort
- ☒ D S is strict

Ans (C)

T_1	T_2
	$\delta(x)$
$\delta(x)$	
	$\delta(y)$
<u>$w(x)$</u> $\delta(y)$	
	<u>$w(x)$</u>
a_1	
	a_2

Recoverable
Cascades

BUT NOT STRICT RECOVERABLE

Q.2



Let S be the following schedule of operations of three transactions T_1 , T_2 and T_3 in a relational database system:

$R_2(Y)$, $R_1(X)$, $R_3(Z)$, $R_1(Y)$, $W_1(X)$, $R_2(Z)$, $W_2(Y)$, $R_3(X)$, $W_3(Z)$

Consider the statements P and Q below:

True \leftarrow P : S is conflict-serializable.

False \leftarrow Q : If T_3 commits before T_1 finishes, then S is recoverable.

Which one of the following choices is correct?

[MCQ: 2021-2M]

☐ A Both P and Q are true.

☒ B P is true and Q is false.


☐ C P is false and Q is true.

☐ D Both P and Q are false.

Ans (B)

T_1	T_2	T_3
	<u>$R(y)$</u>	
<u>$R(x)$</u>		<u>$R(z)$</u>
<u>$R(y)$</u> <u>$W(x)$</u>		
	<u>$R(z)$</u> <u>$W(y)$</u>	
		<u>$R(x)$</u> <u>$W(z)$</u> Commit

Non Recoverable



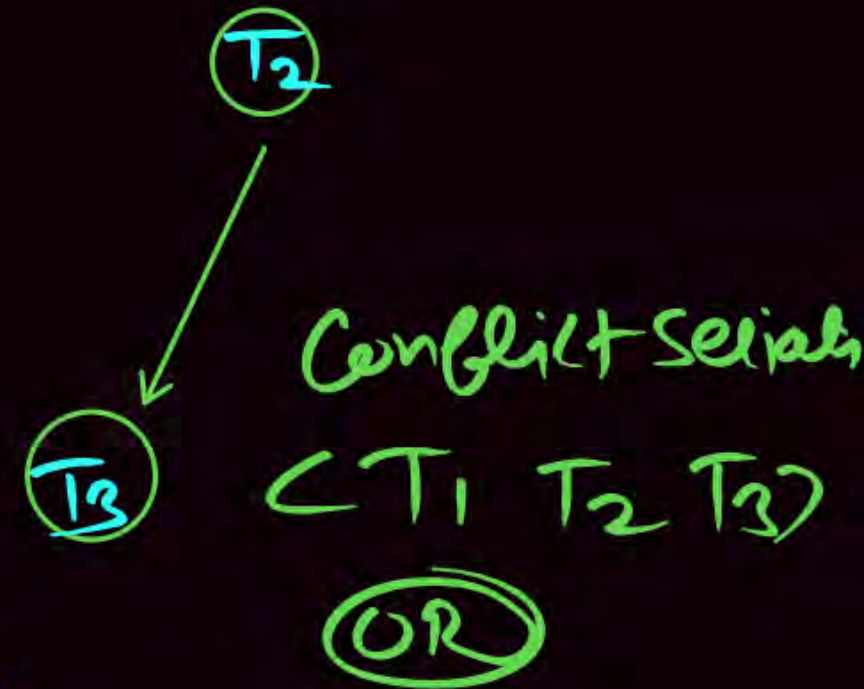
Conflict Serializability
 $\langle T_1, T_2, T_3 \rangle$

OR

Topological Sorting

T_1 : Indegree = 0

T_1	T_2	T_3
	<u>$R(y)$</u>	
<u>$R(x)$</u>		<u>$R(z)$</u>
<u>$R(y)$</u> <u>$W(x)$</u>		
	<u>$R(z)$</u> <u>$W(y)$</u>	
		<u>$R(x)$</u> <u>$W(z)$</u>



By Topological Sorting

Indegree = 0 (T_1)

Serializability order: $\langle T_1, T_2, T_3 \rangle$



Consider a simple checkpointing protocol and the following set of operations in the log.

(start, T4); (write, T4, y, 2, 3); (start, T1);

(commit, T4); (write, T1, z, 5, 7);

(checkpoint);

(start, T2); (write, T2, x, 1, 9); (commit, T2);

(start, T3); (write, T3, z, 7, 2);

If a crash happens now and the system tries to recover using both undo and redo operations, what are the contents of the undo list and the redo list

[MCQ: 2015-2M]

- ☒ A Undo: T3, T1; Redo: T2
- ☐ B Undo: T3, T1; Redo: T2, T4
- ☐ C Undo: none; Redo: T2, T4, T3, T1
- ☐ D Undo: T3, T1, T4; Redo: T2

Redo: T₂

Undo: T₁, T₃

Ans (A).

CHECK POINT : Those Transaction Commit before
Checkpoint Neither Require Redo (no)
Undo operation.

But Those Transaction Commit After Checkpoint
Require Redo operation.

& Those Not Committed till Now Require Undo operation.

UNDO & Redo
Concept

(Write Transaction No. old value new value)
 T_i Data Item $O.V$ $N.V$]

< Write T_9 A 5 7 >

< Write T_{10} A 7 11 >

UNDO: OLD Value
 REDO: New Value

$A = \cancel{11} \cancel{7} 5 \Rightarrow$ A = 5
UNDO CASE
 OLD value

up
UNDO
Bottom

Top
REDO
Bottom

Redo
 new value
 $A = \cancel{5} \cancel{7} 11$
 $A = 11$

CHAPTER 2 : Transaction

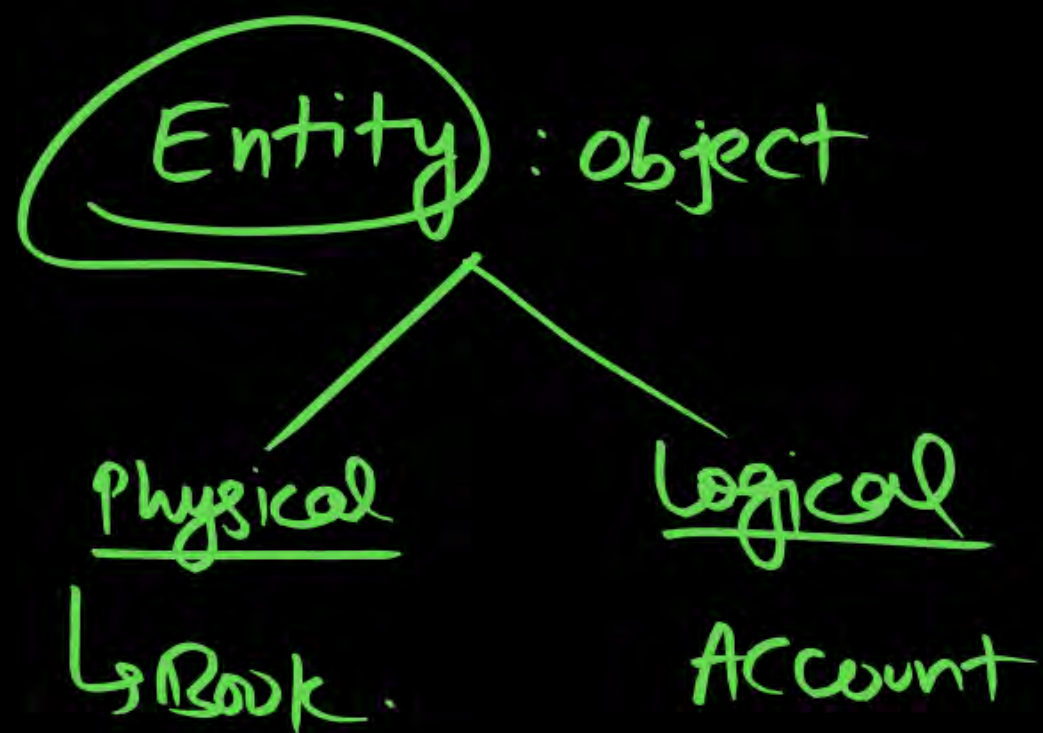
① cc with ^{Chamka 4} Enjoying

② cc

③ c


④ Doubt

Entity – Relationship Model (Conceptual Design)



Relationship ^{→ 'verb.'}

Relationship Set



Diamond



STUDENT

Entity Set : Collection of similar Entity.

Entity Set

STUDENT (Name, age, Gender)

Entity → Abhay 21 Male

Mohita 22 Female

Ajay 25 Male

Entity



An entity is an object that exists and is distinguishable from other objects.

❖ **Example:** Specific person

Entity Set

- ❑ An entity set is a set of entities of the same type that share the same properties.
 - ❖ **Example:** set of all persons, companies

- ❑ Entities have attributes
 - ❖ **Example:** people have names and addresses



Entity sets instructor and student

instructor_ID instructor_name

76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

instructor

Korth

advisor

student_ID student_name

98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

student

Entity Set




It is a set of entities of the same type denoted by a rectangular box in ER diagram. Entity can be identified by a list of attributes which are placed in ovals.

Represented By:



Example:



Relationship Sets

- A relationship is an association among several entities

❖ Example:

44553(Peltier) advisor 22222(Einstein)

student entity relationship set instructor entity

- A relationship set is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where (e_1, e_2, \dots, e_n) is a relationship

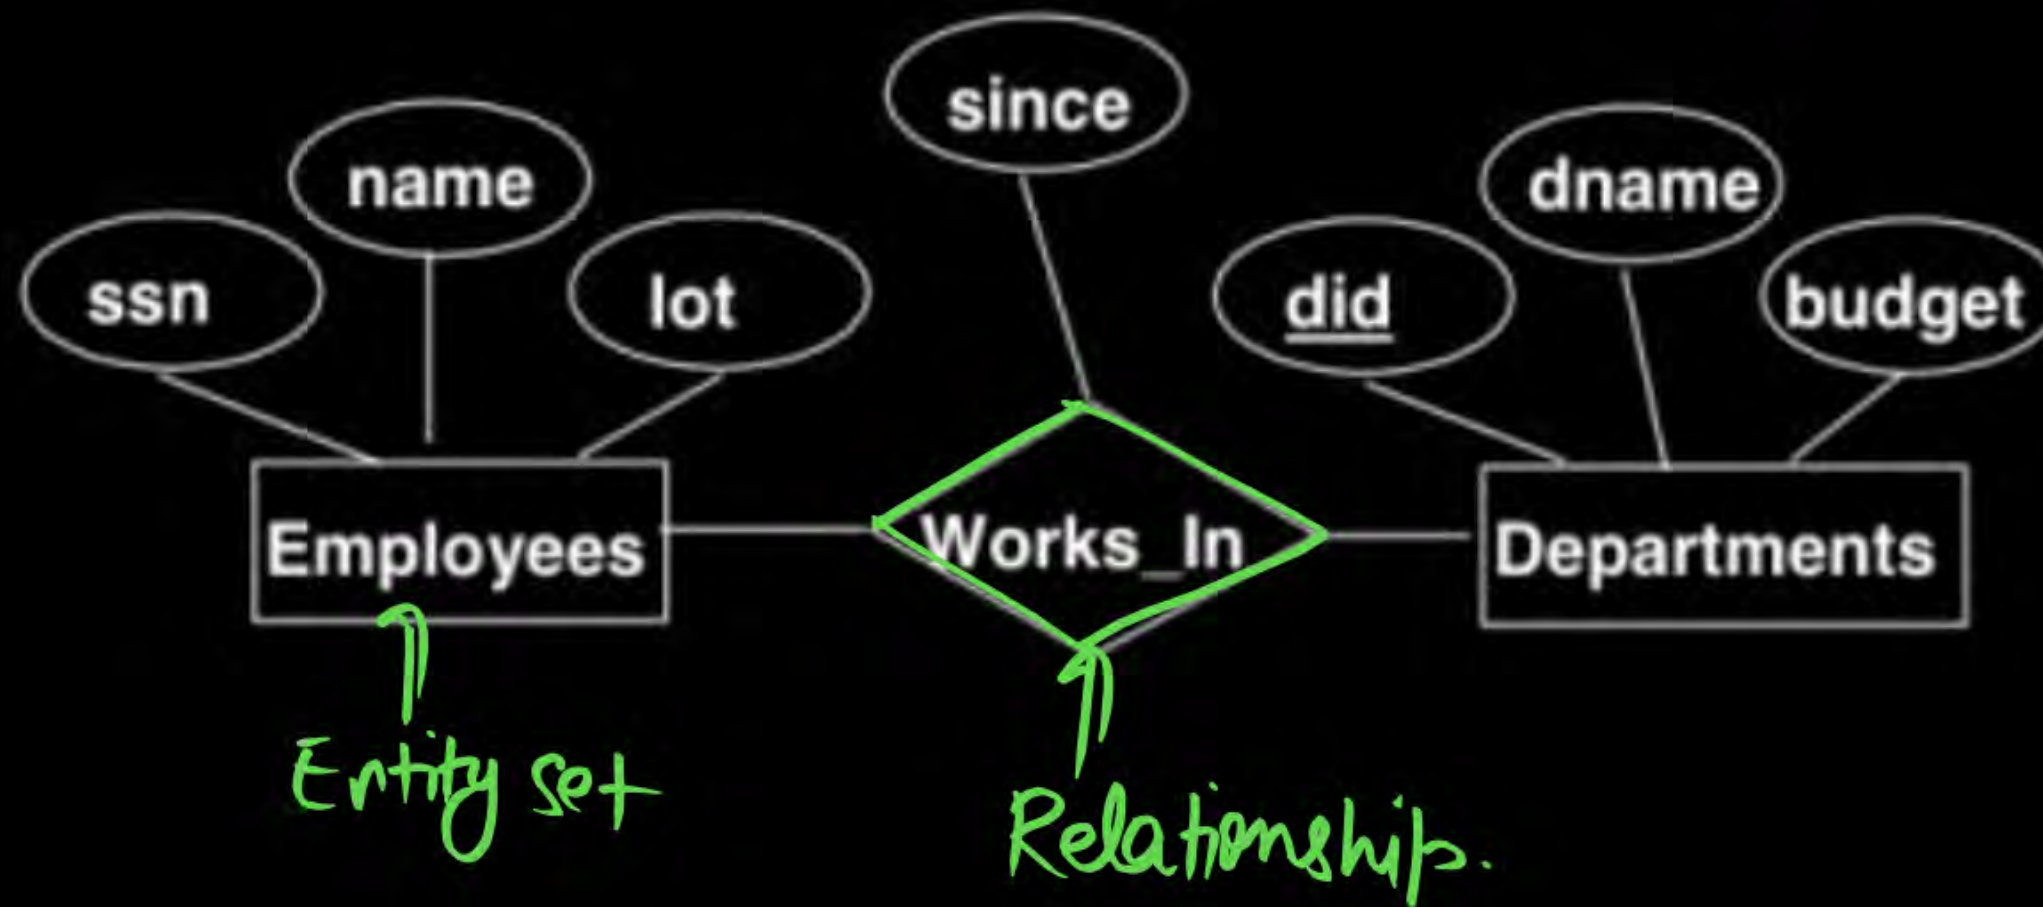

❖ Example:

$(44553, 22222) \in \text{advisor}$

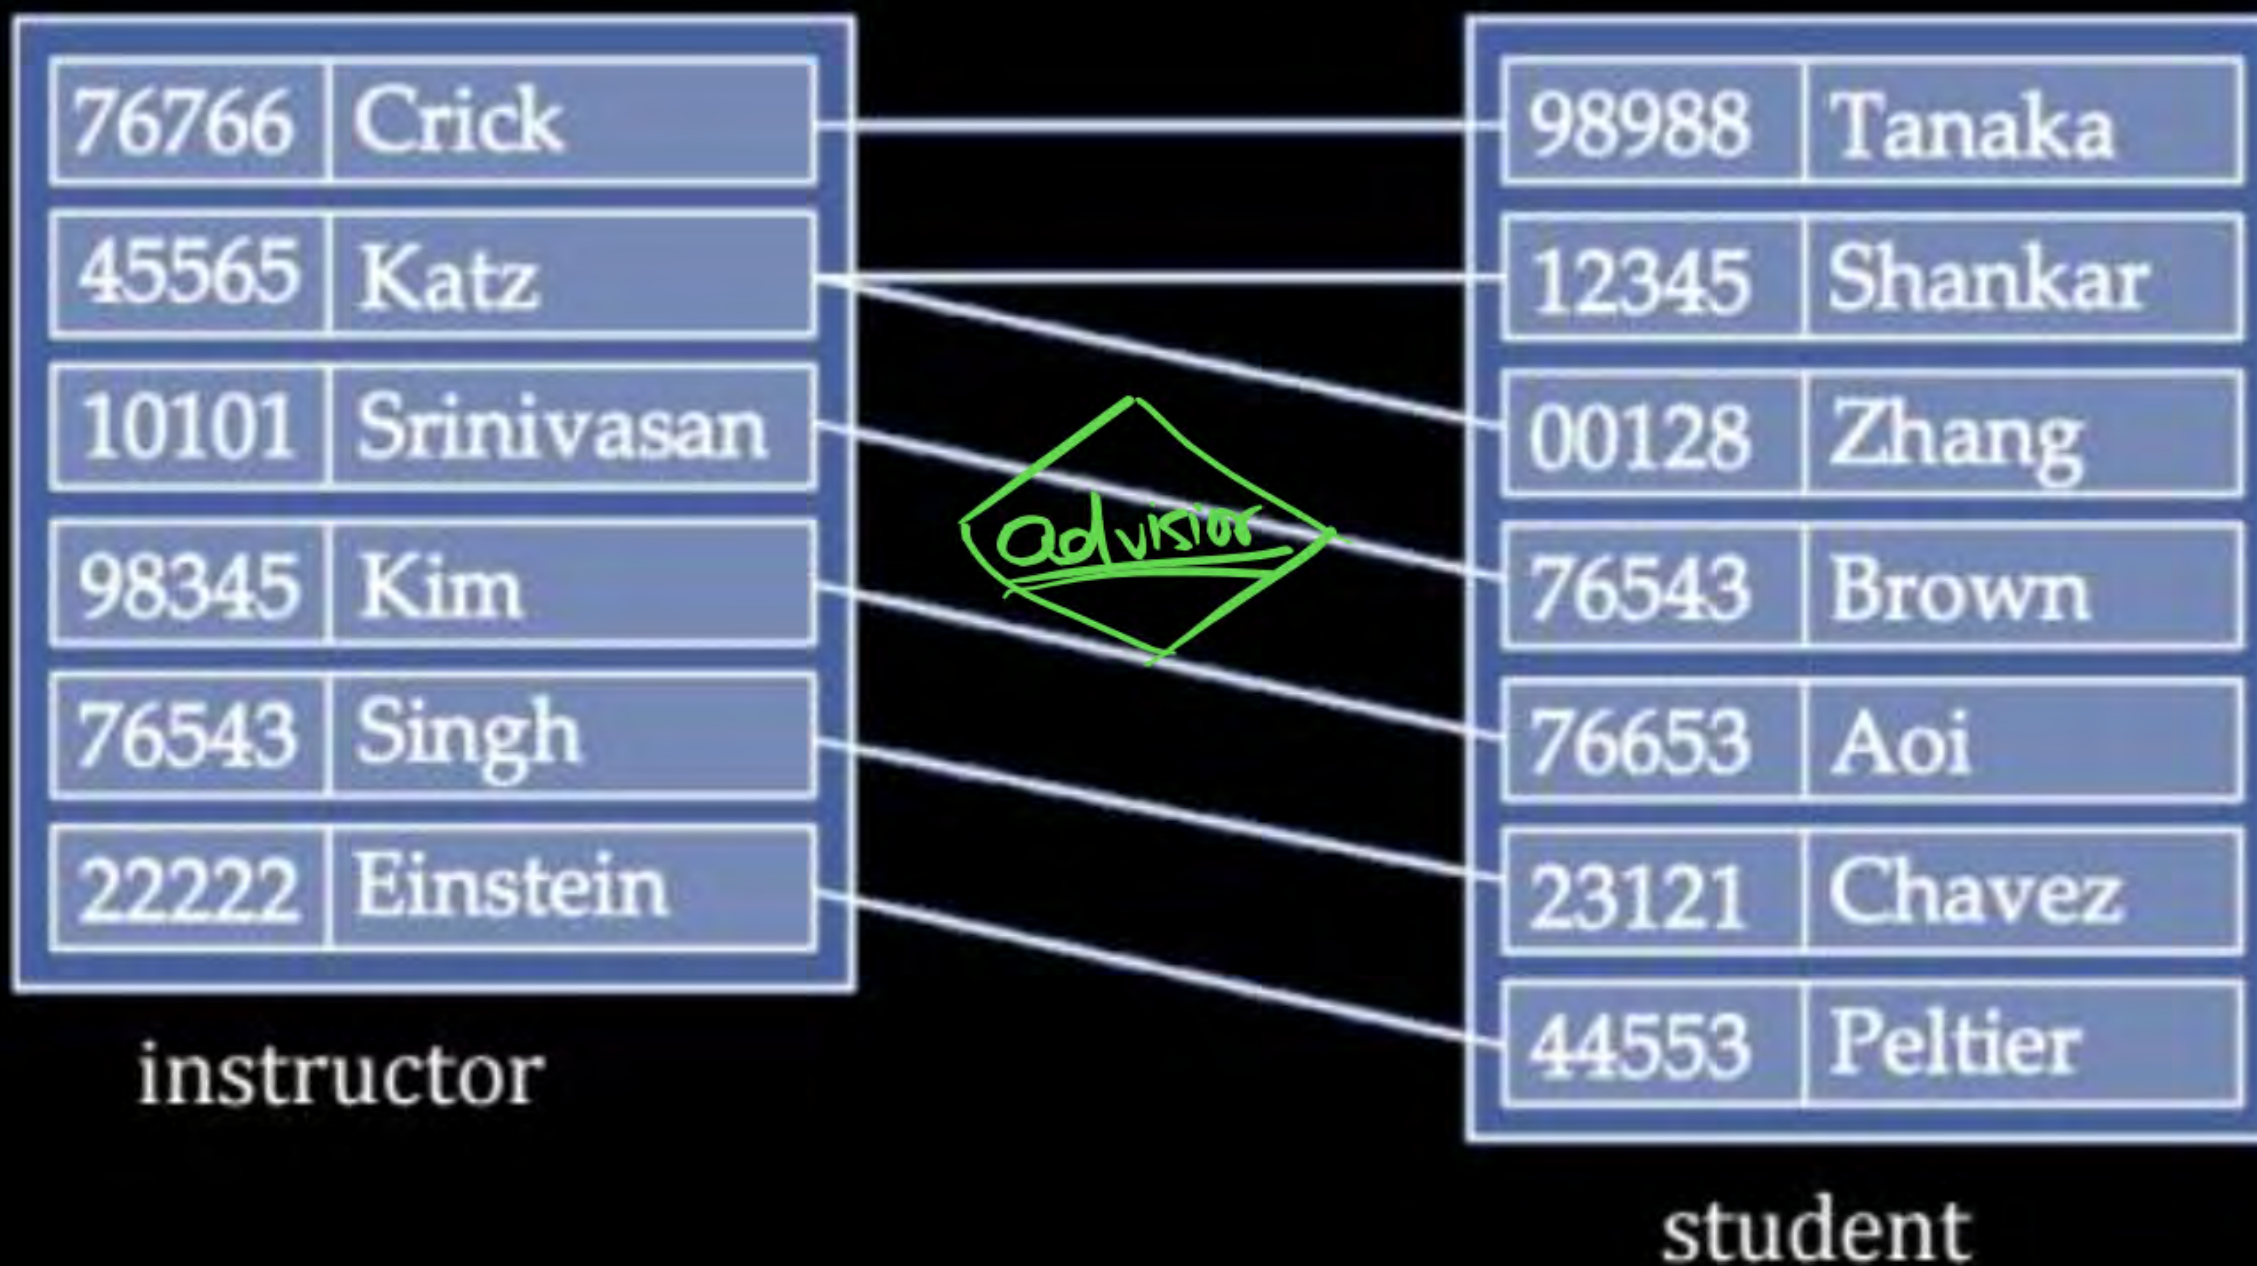
or

Relationship Sets

- Collection of similar relationships.



Relationship Set advisor




①

Entity



Entity Set



②

Relationship



Relationship Set



③ Attribute :



Which describe the Entity.

ER MODEL



Conceptual
Design.



R-MODEL


RDBMS

Relational

Relationship Sets



- ❑ An attribute can also be property of a relationship set.
- ❑ For instance, the advisor relationship set between entity sets instructor and student may have the attribute date which tracks when the student started being associated with the advisor




Attributes



Attributes are properties used to describe an entity.

Attribute types:

- (1) Simple and composite attributes.
- (2) Single-valued and multivalued attributes
- (3) Stored and Derived attributes
- (4) key attribute



① Simple Attribute


Which Can not be Divide further.

eg RollNo
Gender

② Composite Attribute

Which Can be Divide further.

eg




1) Simple & Composite attribute

Simple

Each entity has a single atomic value for the attribute. For example, SSN or Sex (Gender). Address Card Number,

Composite

The attribute may be composed of several components. For example



2) Single-valued and multivalued attributes



Single Valued Attribute: Which

take One Value Per entity.

② Roll No
Gender.

Multivalued Attribute: Which takes
More than One

value Per entity.

②

Mobile

Address

Note

Multivalued Attribute
Represented by



3) Stored and Derived attributes



Stored Attribute: Which does not require any updation.

eg ① D.O.B (Date of Birth)

Derived Attribute: the value is derived from other attribute

eg Age

IB D.O.B : 18/04/2000

then

In 2023 : you are 23 year old

In 2050 : 50 year old

In 2090 : 90 year old


In 2110 : 110 year old.

4) Key and Descriptive attributes

Key Attribute:

Which uniquely Identify an entity in the entity set.

Roll Number



Descriptive Attribute:

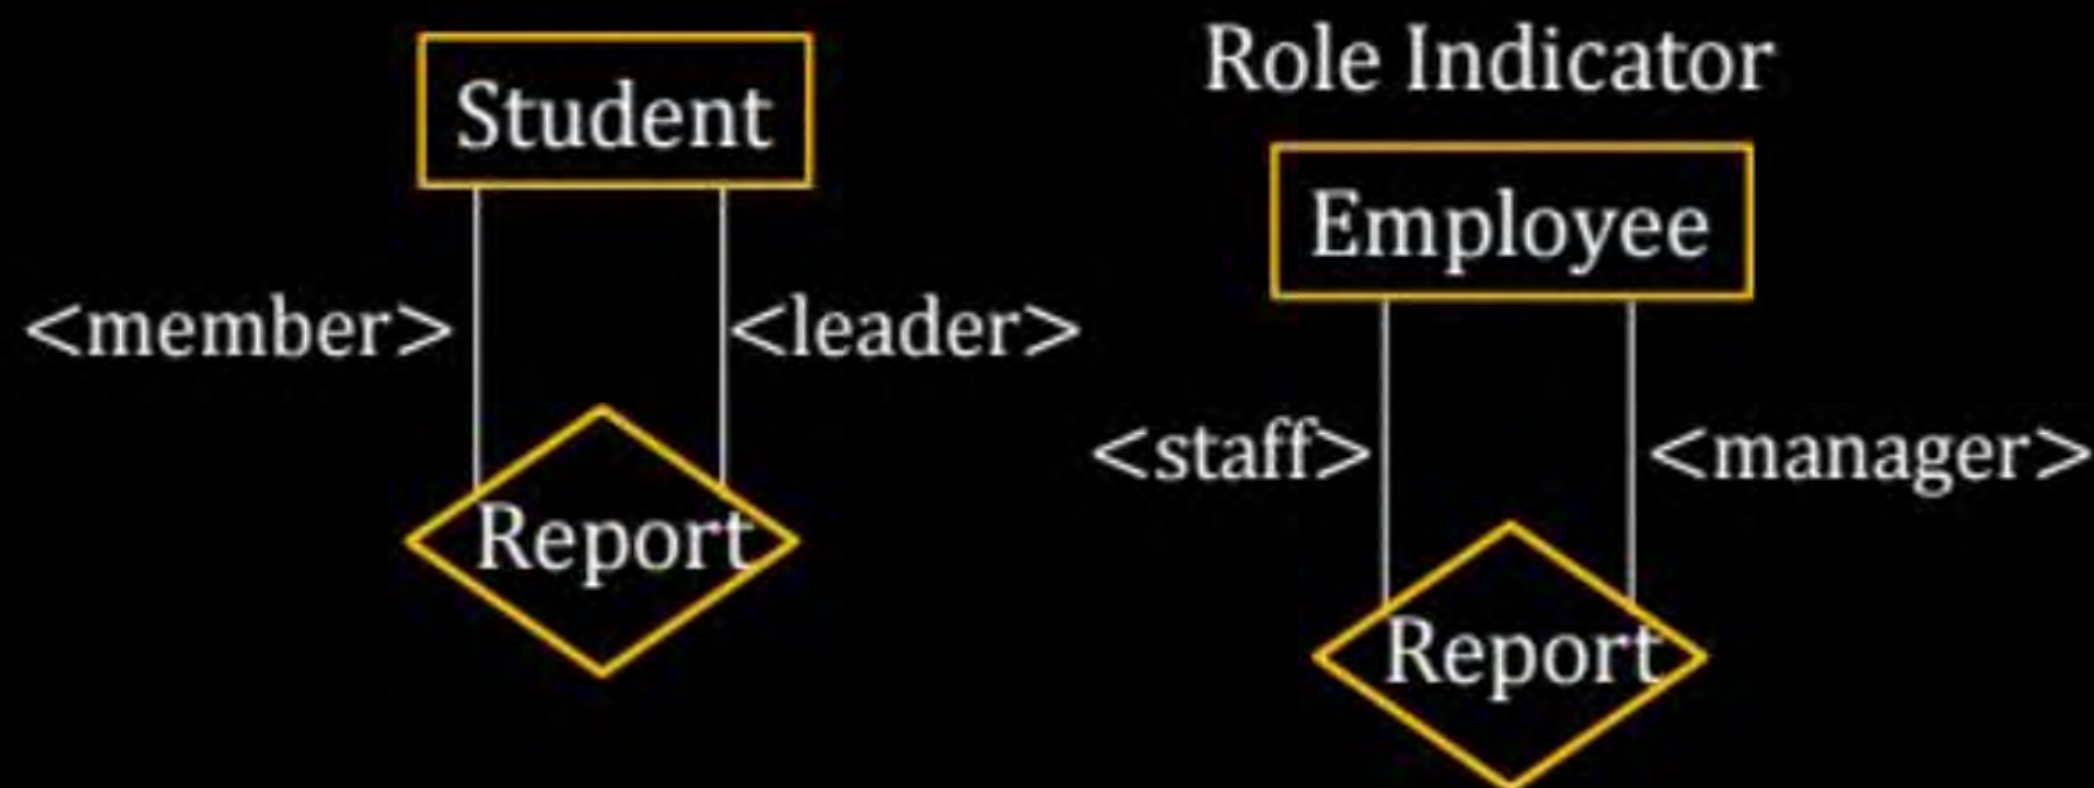
Which gives information about the relationship set



Degree of Relationship Set

Degree of Relationship Set: Specifies the numbers of Entity set participate in a relationship set

1) UNARY: Relationship among two entities of the same entity set
[Recursive Relationship Set]




2) Binary Relationship: The relationship among two entity set



2) Ternary Relationship: The relationship among three entity set



3) n-ary Relationship:
The relationship among n-entity set





Participation Constraint

If every entity in the entity set participates in a relationship set is called

Total participation denoted by double line (thick line)
otherwise it is called partial participation (thin line or single line)

Q.

Each department is managed by at least one employee



Participation



- If every entities of entity set are participated with relationship set then it is total participation (100% participation) otherwise it will be partial participation ($< 100\%$ participation)

Example : Consider Emp and Dept entity set.


Manages relationship set such that each dept must have manager.




Mapping Cardinality Constraints

- ❑ Express the number of entities to which another entity can be associated via a relationship set.
- ❑ Most useful in describing binary relationship sets.
- ❑ For a binary relationship set the mapping cardinality must be one of the following types:
 - ❖ One to one
 - ❖ One to many
 - ❖ Many to one
 - ❖ Many to many

Mapping Cardinalities




One to one



One to many


Note: Some elements in A and B may not be mapped to any elements in the other set

Mapping Cardinalities



(a)

Many to
one



(b)

Many to many

Note: Some elements in A and B may not be mapped to any elements in the other set


Mapping [Cardinality constraints of relationship set]




One mapping : At most one (0 or 1)

Many mapping : 0 or more (0 *)


Binary Relationship Mapping (One : One)



Binary Relationship Mapping (One : Many)




Binary Relationship Mapping (Many to One)



Mapping [Cardinality constraints of relationship set]



Binary Relationship Mapping (Many to Many)



Any Doubt ?



**THANK
YOU!**

