

COMPUTER SCIENCE



Database Management System

Transaction & Concurrency Control

Lock Based Protocol - Part - 03

Lecture_10

Vijay Agarwal sir



An orange diamond-shaped sign with a black border and the text 'TOPICS TO BE COVERED' in black capital letters.

TOPICS
TO BE
COVERED

A red diamond-shaped marker with a white border and the number '01' in white.

01

Lock Based Protocol

A red diamond-shaped marker with a white border and the number '02' in white.

02

Time Stamp Protocol



LOCK BASED Protocol

SHARED LOCK [S] \Rightarrow Only Read

Exclusive lock [X] \Rightarrow write / write
Read

Some Data Item		T_i	
		S	X
T_j	S	YES	NO
	X	No	No



Consider the following database schedule with two transactions, T_1 and T_2 .

$S = r_2(X); r_1(X); r_2(Y); w_1(X); r_1(Y); w_2(X); a_1; a_2$

where $r_i(Z)$ denotes a read operation by transaction T_i on a variable Z , $w_i(Z)$ denotes a write operation by T_i on a variable Z and a_i denotes an abort by transaction T_i

Which one of the following statements about the above schedule is TRUE?

[MCQ:2016–2M]

- A** S is non-recoverable
- B** S is recoverable, but has a cascading abort
- C** S does not have a cascading abort
- D** S is strict

Q.2



Let S be the following schedule of operations of three transactions T_1 , T_2 and T_3 in a relational database system:

$R_2(Y), R_1(X), R_3(Z), R_1(Y), W_1(X), R_2(Z), W_2(Y), R_3(X), W_3(Z)$

Consider the statements P and Q below:


P : S is conflict-serializable.

Q : If T_3 commits before T_1 finishes, then S is recoverable.

Which one of the following choices is correct?

[MCQ: 2021-2M]

- A** Both P and Q are true.
- B** P is true and Q is false.
- C** P is false and Q is true.
- D** Both P and Q are false.



Consider a simple checkpointing protocol and the following set of operations in the log.

(start, T4); (write, T4, y, 2, 3); (start, T1);

(commit, T4); (write, T1, z, 5, 7);

(checkpoint);

(start, T2); (write, T2, x, 1, 9); (commit, T2);

(start, T3); (write, T3, z, 7, 2);

If a crash happens now and the system tries to recover using both undo and redo operations, what are the contents of the undo list and the redo list

[MCQ: 2015–2M]

- A** Undo: T3, T1; Redo: T2
- B** Undo: T3, T1; Redo: T2, T4
- C** Undo: none; Redo: T2, T4, T3, T1
- D** Undo: T3, T1, T4; Redo: T2



Lock - S(A)

Read(A)

unlock - S(A)

Lock - X(A)

Write(A) \otimes Write(A)
Read(A)

Unlock - X(A).

Lock-Based Protocols

- ❑ A lock is a mechanism to control concurrent access to a data item
- ❑ Data items can be locked in two modes:
 1. exclusive (X) mode. Data item can be both reads as well as written. X-lock is requested using **lock-X** instruction.
 2. Shared (S) mode. Data item can only be read. S-lock is requested using **lock-S** instruction.
- ❑ Lock requests are made to concurrency-control manager. Transaction can proceed only after request is granted.

Lock-Based Protocols (Cont.)

❑ Lock-compatibility matrix

	S	X
S	true	false
X	false	false

- ❑ A transaction may be granted a lock on an item if the requested lock is compatible with lock already held on the item by other transactions
- ❑ Any number of transactions can hold shared locks on an item
- ❑ But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.

Schedule with Lock Grants

- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols enforce serializability by restricting the set of possible schedules.

T ₁	T ₂	Concurrency-control manager
lock-X(B)		grant-X(B, T ₁)
read(B)		
B:=B - 50		
write(B)		
unlock(B)		
	lock-S(A, T ₂)	grant-S(A, T ₂)
	read(A)	
	unlock(A)	
	lock-S(B)	grant-S(B, T ₂)
	read(B)	
	unlock(B)	
	display(A+B)	
lock-X(A)		grant-X(A, T ₁)
read(A)		
A:=A + 50		
write(A)		
unlock(A)		

2 Phase locking Protocol [2PL]

① Growing Phase [Acquire Lock]

② Shrinking Phase [Release Lock]

Here we Issue Lock & Unlock Request into the Two Phases.

① Growing Phase : In Growing Phase Transaction
May obtain the lock But
Must Not Release Any Lock.

② Shrinking Phase : In Shrinking Phase
Transaction Release the lock
But Must Not Ask for Any New Lock.

Each Transaction finish first its Growing phase
then Shrinking phase.

The Two-Phase Locking Protocol

❑ A protocol which ensures conflict-serializable schedules.

❑ Phase 1: **Growing Phase**

❖ Transaction may obtain locks


❖ Transaction may not release lock

❑ Phase 2: **Shrinking Phase**

❖ Transaction may release locks

❖ Transaction may not obtain locks

❑ The protocol assures serializability. It can be proved that transactions can be serialized in the order of their lock points (i.e., the point where a transaction acquired its final lock).




Serializability order (Equivalent Serial Schedule) is Determined by Lock POINT.

Lock POINT

→ Position of last Lock (or) first Unlock operation.

is a point FROM WHERE Shrinking Phase of the transaction Starts.



Serializability : $\langle T_2 T_3 T_1 \rangle$
order

2PL

①

Read(A)
A = A - 100
Write(A)

Read(B)
B = B + 100
Write(B)

T₁

T₂

Read(A)
Read(C)

T₁

T₂

Lock-X(A)
Read(A)
A = A - 100
Write(A)

Lock-X(B)

Unlock-X(A)

Lock(S)
Point

Read(B)
B = B + 100
Write(B)
Unlock-X(B)

Lock-S(A)

Lock-S(A)
Read(A)

Lock-S(C)

Unlock-S(A)
Read(C)
Unlock-S(C)

Concurrency control manager
Grant-X(A, T₁)

X Not allowed

Grant-X(B, T₁)

Revoke-X(A, T₁)

Grant-S(A, T₂)

Grant-S(B, T₂)

T₁ → T₂

2PL

	T_1	T_2
②	Read(A) $A = A - 100$ <u>Write(A)</u>	<u>Read(A)</u> Commit

Commit |

Irrecoverable / Schedule
 Non recoverable

T_1	T_2
Lock-X(A) Read(A) $A = A - 100$ write(A) unlock-X(A)	Lock-S(A) Read(A) unlock-S(A) Commit
Commit	

Irrecoverable schedule But followed by 2PL.

③

T ₁	T ₂
<u>R(A)</u>	<u>W(B)</u>
<u>R(B)</u>	<u>W(A)</u>

Denied by T₂ → R(B)

T ₁	T ₂
S(A) R(A)	<u>X(B)</u> W(B)

X(A) Denied by T₁

Deadlock

T ₁	T ₂
W(A)	<u>R(B)</u>
W(B)	<u>R(A)</u>

Denied by T₂ → X(B)

T ₁	T ₂
X(A) W(A)	<u>S(B)</u> R(B)

S(A) Denied by T₁

Deadlock

	T_1	T_2	T_3	T_4	T_5
Denied by T_2 ↳ $X(A)$		<u>$S(A)$</u>			
Denied by T_3 ↳ $X(A)$		unlock- $S(A)$	<u>$S(A)$</u>		
Denied by T_4 ↳ $X(A)$			unlock- $S(A)$	<u>$S(A)$</u>	
Denied by T_5 ↳ $X(A)$				unlock- $S(A)$	<u>$S(A)$</u>
Granted ↳ $X(A)$					unlock- $S(A)$

STARVATION.

Important Points about 2PL.

① 2PL Ensure Conflict Serializable Schedule.
↳ Serializability

OR

If a Schedule is Followed by 2PL then It Ensure Conflict Serializability

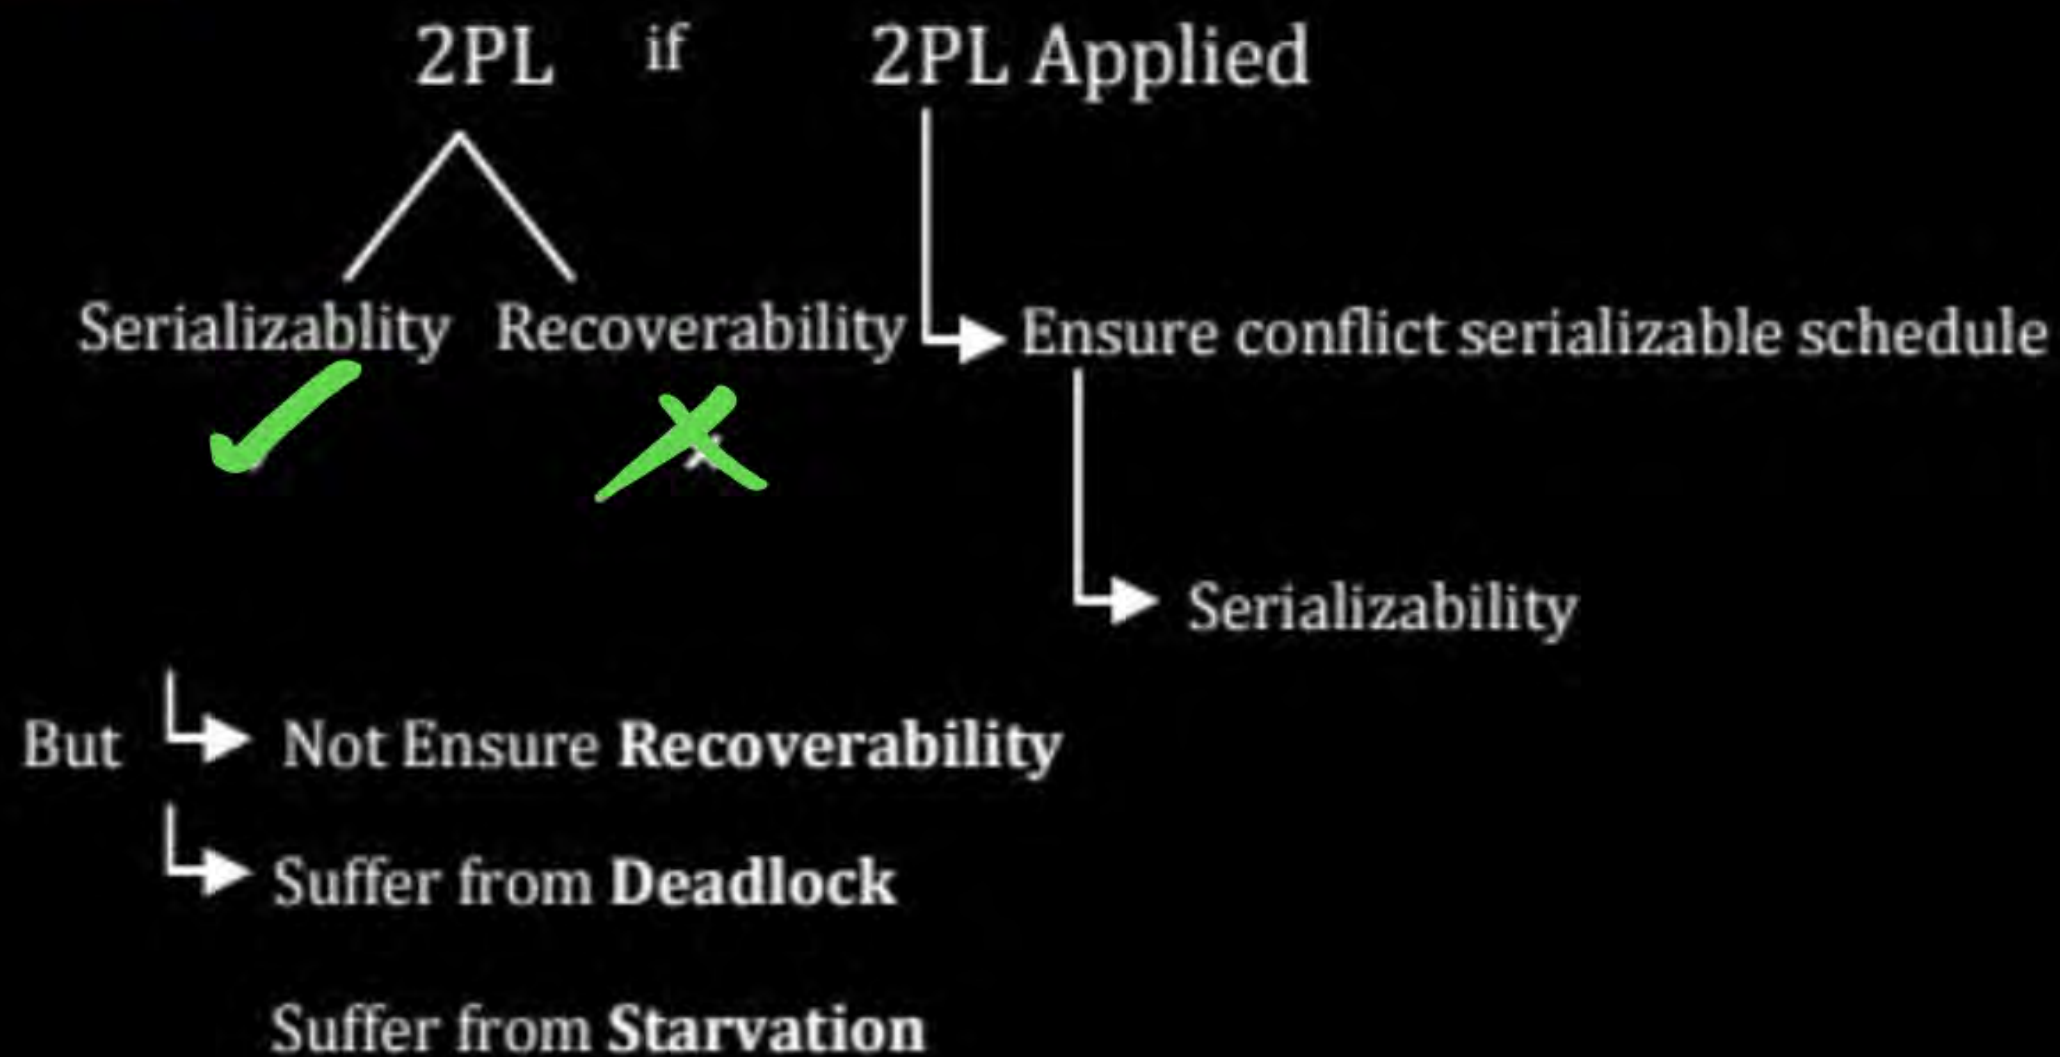
② Serializability order Determined by LOCK POINT.

③ 2PL Not Ensure Recoverability.

④ 2PL Not Free From Deadlock | Subbel from Deadlock.

⑤ 2PL Subbel from Starvation | Not Free From Starvation.

2PL



2PL

Serializable
⇓
Conflict
Serializable

But
Recoverable
Not Ensure.

② STRICT 2PL:

(2PL) + All Exclusive Lock Taken by the Transaction until Commit/Rollback.

(OR)

2PL + All Exclusive Lock Release After Commit/Rollback of the transaction.

②

S(A)
X(B)
X(C)
S(D)
=

Unlock(A)
Unlock(D)

Commit

Unlock - X(B)

Unlock - X(C)

T ₁	T ₂
X(A)	
W(A)	
<u>C/R</u>	
Unlock - X(A)	
	S(A) / X(A)
	R(A) / W(A)

Strict Recoverable

①

T_1	T_2
$w(A)$	
	$R(A)$
clr	
	• Commit

Recoverable

②

T_1	T_2
$w(A)$	
clr	
	$R(A)$

Cascadeless

T_1	T_2
$w(A)$	
clr	
	$R(A) / w(A)$

Strict Recoverable

STRICT 2PL \Rightarrow 2PL +

- Conflict serializable.
- ↳ Ensure
 - Recoverable Schedule
 - Cascades
 - Strict Recoverable
 - No Cascading Rollback.

But Subber / Not Free From

- ↳ Deadlock
- ↳ Starvation.

③ Rigorous 2PL

2PL + All locks (SHARED[S] & Exclusive[X]) Held until Commit/Rollback of the transaction **OR** Release After Commit/Rollback.

T_i

S(A)

X(B)

X(C)

S(D)

CR

U(A)

U(B)

U(C)

U(D)

Rigorous 2PL are $\begin{cases} \rightarrow 2PL \\ \rightarrow \text{Strict 2PL} \end{cases}$

But Suffer from


- Deadlock
- Starvation.

The Two-Phase Locking Protocol (Cont.)



- ❑ Two-phase locking does not ensure freedom from deadlocks
- ❑ Extensions to basic two-phase locking needed to ensure recoverability of freedom from cascading roll-back
 - ❖ ① Strict two-phase locking: a transaction must hold all its exclusive locks till it commits/aborts.
 - Ensures recoverability and avoids cascading roll-backs *→ Cascades, Strict Recoverable.*
 - ❖ ② Rigorous two-phase locking: a transaction must hold all locks till commit/abort. *→ (SHARED & Exclusive)*
 - Transactions can be serialized in the order in which they commit.
- ❑ Most databases implement rigorous two-phase locking, but refer to it as simply two-phase locking

④ Conservative 2PL : Which State Acquire(taken)
all the Locks before transaction
starts (at the beginning) & Release
ALL the Locks After Commit / Rollback.



Conservative → 2PL
 ↳ Strict 2PL

No Deadlock

Only suffer from starvation.

1	2	3	4	5
Lock-S(A) R(A) Lock-X(B) R(B) Unlock(A) W(B) Unlock(B) Commit	Lock-S(A) R(A) Lock-X(B) Unlock(A) R(B) W(B) Commit Unlock(B)	Lock-S(A) R(A) Lock-X(B) R(B) W(B) Commit Unlock(A) Unlock(B)	Lock-S(A) Lock-X(B) R(A) R(B) W(B) Commit Unlock(A) Unlock(B)	Lock-S(A) R(A) Unlock(A) Lock-X(B) R(B) W(B) Unlock(B) Commit

2PL only

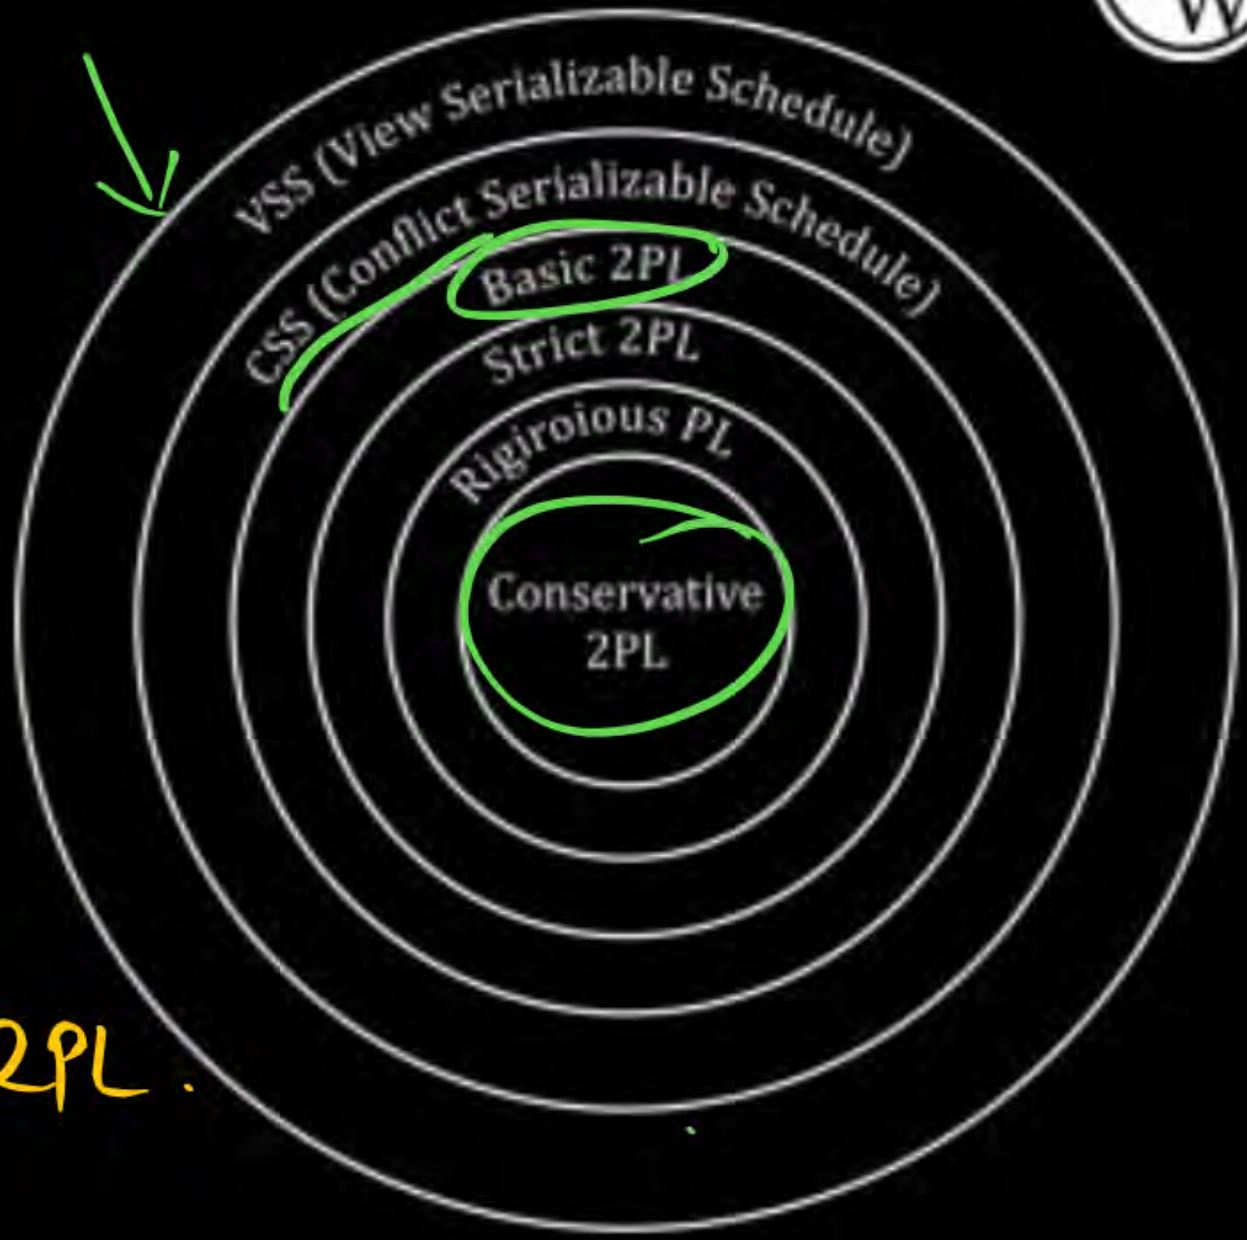
Strict 2PL
2PL

Rigorous 2PL
Strict 2PL
2PL

Conservative 2PL
Rigorous 2PL
strict 2PL
2PL

T_i	T_j
R(A)	
	w(p)
R(B) w(B)	

Not 2PL






TIMESTAMP BASED CONCURRENCY CONTROL

Timestamp Based Protocol

Unique Time Stamp value is assigned to each Transaction,
When Transaction Enter into the System

$TS(T_1) = 10$
 $TS(T_2) = 30$
 $TS(T_3) = 20$
 $TS(T_4) = 40$



Timestamp-Based Protocols

- ❑ Each transaction T_i is issued a timestamp $TS(T_i)$ when it enters the system.
 - ❖ Each transaction has a unique timestamp
 - ❖ Newer transaction have timestamp strictly greater than earlier ones
 - ❖ Timestamp could be based on a logical counter
 - Real time may not be unique
 - Can use (wall-clock time, logical counter) to ensure
- ❑ Timestamp-based protocols manage concurrent execution such that
time-stamp order = serializability order
- ❑ Several alternative protocols based on timestamps

If Transaction T_j enter, After the transaction T_i

then

$$TS(T_j) > TS(T_i)$$

Time Stamp of T_j

$$TS(T_j) > TS(T_i)$$

T_i
✓
OLDER
Transaction

T_j
↓
younger (Newest)
Transaction

Serializability $\langle T_i T_j \rangle$
order

T_i followed by
 T_j

There are 2 Type of Time Stamp.

① ① Transaction Time Stamp $TS(T_i)$: Fixed

② ② Data Item Time Stamp: $\begin{cases} RTS(A) \\ WTS(A) \end{cases} \left\{ \begin{array}{l} \text{variable} \end{array} \right\}$

$RTS(A)$: Read Time Stamp
on Data Item A

$WTS(A)$: Write Time Stamp on
Data Item A.

① Read-TimeStamp (RTS(A))

$\begin{pmatrix} 10 \\ T_1 \end{pmatrix}$	$\begin{pmatrix} 20 \\ T_2 \end{pmatrix}$	$\begin{pmatrix} 30 \\ T_3 \end{pmatrix}$
R(A)	R(A)	R(A)

Initially

RTS(A): 30

Denote the Highest Transaction TimeStamp which perform R(A) operation successfully.

$$RTS(A) = 0;$$

$$RTS(A) = 10; \max(0, 10)$$

$$RTS(A) = 20; \max(10, 20)$$

$$RTS(A) = 30; \max(20, 30)$$

② Write Time Stamp (WTS(A))

Denote the Highest Transaction Time Stamp that perform write(A) operation successfully.

10 T_1	20 T_2	30 T_3
$W(A)$	$W(A)$	$W(A)$

$$WTS(A) = 30$$

$$WTS(A) = 0 \quad (\text{initially})$$

$$WTS(A) = 10 \quad \max(0, 10)$$

$$WTS(A) = 20 \quad \max(10, 20)$$

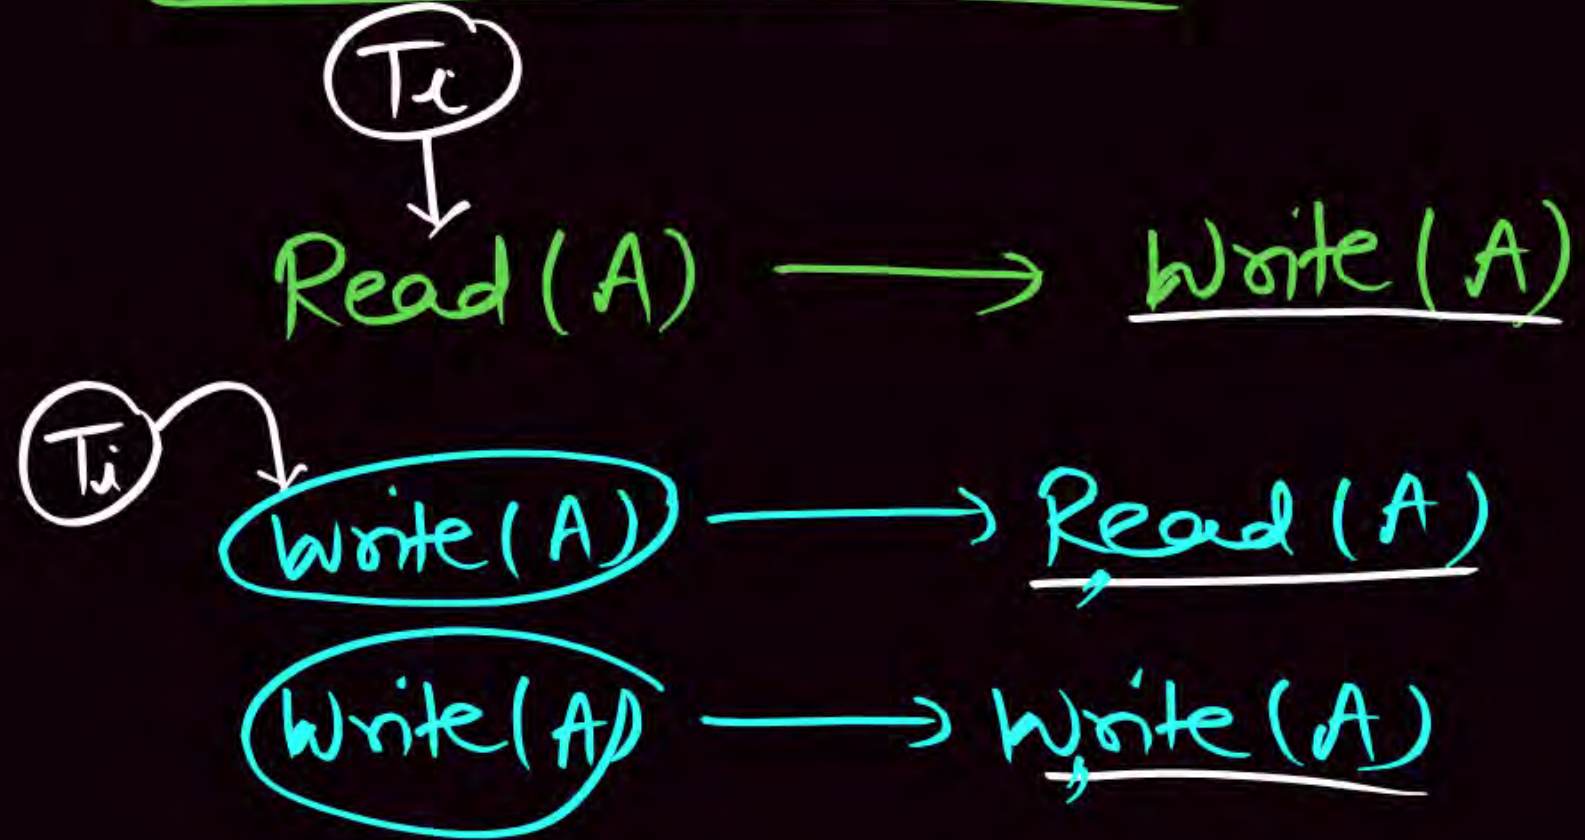
$$WTS(A) = 30 \quad \max(20, 30)$$

Timestamp-Based Protocols

The timestamp ordering (TSO) protocol

- ❑ Maintains for each data Q two timestamp values:
 - ❖ W-timestamp(Q) is the largest time-stamp of any transaction that executed write (Q) successfully.
 - ❖ R-timestamp (Q) is the largest time-stamp of any transaction that executed read (Q) successfully.
- ❑ Imposes rules on read and write operations to ensure that
 - ❖ any conflicting operations are executed in timestamp order
 - ❖ out of order operations cause transaction rollback

Conflict operation



T_i \Rightarrow Read(Q)

If $TS(T_i) < \underline{WTS(Q)}$; Reject

T_i : Write(Q)

$TS(T_i) < RTS(Q)$; Reject

$TS(T_i) < WTS(Q)$; Reject

I: T_i – Read(Q) (Transaction T_i Issue R(Q) Operation)

- (i) If $TS(T_i) < \underline{WTS(Q)}$: Read operation Reject & T_i Rollback.
- (ii) If $TS(T_i) \geq WTS(Q)$: Read operation is allowed
and Set Read – $TS(Q)$ = $\max[RTS(Q), TS(T_i)]$

II: T_i – Write(Q) (Transaction T_i Issue Write(Q) Operation)

- (i) If $TS(T_i) < RTS(Q)$: Write operation Reject & T_i Rollback.
- (ii) If $TS(T_i) < WTS(Q)$: Write operation Reject & T_i Rollback.
- (iii) Otherwise execute write (Q) operation
Set Read $WTS(Q) = TS(T_i)$

Read



$T_i \equiv T_2$
 \Downarrow
 $TS(T_2) = 20$



$W(Q)$



$$WTS(Q) = 30$$

$$TS(T_i) = TS(T_2) = \textcircled{20}$$

Write (Q) $T_i = T_2 \Rightarrow TS(T_2) = 2021$

$$TS(T_i) < RTS(Q)$$

year.

<u>2020</u> T_1	<u>2021</u> T_2	<u>2023</u> T_3
	<u>$W(Q)$</u>	
		$R(Q)$

$$RTS(Q) = \underline{2023}$$

$$TS(T_2) = \underline{2021}$$

$$TS(T_i) < WTS(Q)$$

<u>2020</u> T_1	<u>2021</u> T_2	<u>2023</u> T_3
	$W(Q)$	
		$W(Q)$

$$WTS(Q) = \underline{2023}$$

$$TS(T_2) = 2021$$

Timestamp-Based Protocols (Cont.)


- ❑ Suppose a transaction T_i issues a **read** (Q)
 1. If $TS(T_i) \leq W\text{-timestamp}(Q)$, then T_i needs to read a value of Q that was already overwritten.
 - Hence, the **read** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) \geq W\text{-timestamp}(Q)$, then the **read** operation is executed, and $R\text{-timestamp}(Q)$ is set to.

$$\max(R\text{-timestamp}(Q), TS(T_i)).$$


Timestamp-Based Protocols (Cont.)

- Suppose a transaction T_i issues **write**(Q)
 1. If $TS(T_i) < \text{R-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that the value would never be produced.
 - Hence, the **write** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) < \text{W-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q.
 - Hence, this **write** operation is rejected, and T_i is rolled back.
 3. Otherwise, the **write** operation is executed, and $\text{W-timestamp}(Q)$ is set to $TS(T_i)$.


(1)




(2)



(3)



(4)



Thomas' Write Rule

- ❑ Modified version of the timestamp-ordering protocol in which obsolete **write** operations may be ignored under certain circumstances.
- ❑ When T_i attempts to write data item Q , if $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of $\{Q\}$.
 - ❖ Rather than rolling back T_i as the timestamp ordering protocol would have done, this **{write}** operation can be ignored.
- ❑ Otherwise this protocol is the same as the timestamp ordering protocol.
- ❑ Thomas' Write Rule allows greater potential concurrency.
 - ❖ Allows some view-serializable schedules that are not conflict-serializable.

Thomas Write Rule (View Serializability)



1. $TS(T_i) < RTS(Q)$: Rollback
2. $TS(T_i) < WTS(Q)$: Write operation is Ignored and No Roll back


Same as TSP

Time Stamp Protocol: Ensure serializability deadlock free but starvation possible

Deadlock Prevention Algorithm

- (1) Wait-Die
- (2) Wound-wait

Older Younger





**THANK
YOU!**

