

# CS & IT ENGINEERING

**Operating Systems**  
**Miscellaneous Topics**

**Lecture No. 1**



By- Dr. Khaleel Khan Sir





TOPICS TO BE  
COVERED

**Fork() System Call**

**Threads**

**Monitors**



```
main()
```

```
{  
  int a;
```

```
  if (fork() == 0)  
  {
```

// child //

```
    a = a + 5;
```

```
    print(a);  $\Rightarrow$  "4"
```

```
  }  
  else  
  {
```

// Parent //


```
    a = a - 5;
```

```
    print(a);  $\Rightarrow$  "x"
```

```
  }  
}
```

$u \propto x$

- a)  $u = x + 10$
- b)  $x = u + 10$
- c)  $u = x$
- d)  $u = x - 5$ ;





# Q.1

```
main ()
{
    int i, n;
    for (i = 1; i <= n; ++i)
    if (fork () == 0) ;
    print ("*");
}
```

print ("\*"); → outside for-loop  
All : child + Parent

No. of times '\*' gets  
 printed  $(2^n - 1)$  ✓

$2^n$  ✓

# Q. 2

```
main ()
{
    int i, n;
    for (i = 1; i <= n; ++i)
    {
        1. fork ();
        2. print ("*");
    }
}
```

\*\*

H/w : No. of times '\*' \_\_\_\_\_

# Q3

```
for i ← 1 to n
{
    print ('*');
    fork();
}
```

\*



Q. 3

main ()

{

int i, n;

for (i = 1; i &lt;= n; ++ i)

{

print ("\*");

fork ();

}

}



Q. 4

main ()  
{

int a = 1, b = 2, c = 3; *d = 4;*

a += b += ++c;

print (a, b, c); *// Parent //*

if (fork () == 0)

{

int d;

++a; ++b; --c;

print (a, b, c);

if (fork () == 0)

{

d = a + b + c;

print (a, b, c, d);

}

*// child 1*

*// child 2*

else

{

--a; --b;

c = a + b; d = a + b + c;

print (a, b, c, d);

}

}

else

{

c += b += ++a;


print (a, b, c);

}


print (d);

*All Processes*





'd' is NOT defined in Parent





Q. 5

The following C program is executed on a Unix/Linux system:

```
#include <unistd.h>
```

```
int main()
```

```
{   int i;
```

```
    for (i = 0; i < 10; i++)
```

```
        if (i % 2 == 0) fork();
```

```
}
```

*even (5) → '5' Times*

The total number of child processes created is  $\underline{(2^5 - 1) = 31}$  ✓



main()  $\rightarrow$   $c = c + (++a) * (++b);$



```

{
    int a=1, b=2, c=0;

    c += ++a * ++b;
    Print(a, b, c);
    if (fork() == 0)
    {
        a += b += c += 5;
        print(a, b, c);
        c++;
    }
    else
    {
        b += c += a += 3;
        print(a, b, c);
        c++;
    }
}

```

print(c);





# Threads & Multithreading:

→ (A Light weight Process)

→ Part of Process

→ uses Resources  
of Process

→ Context Switch  
ovhd b/w the  
Threads is less

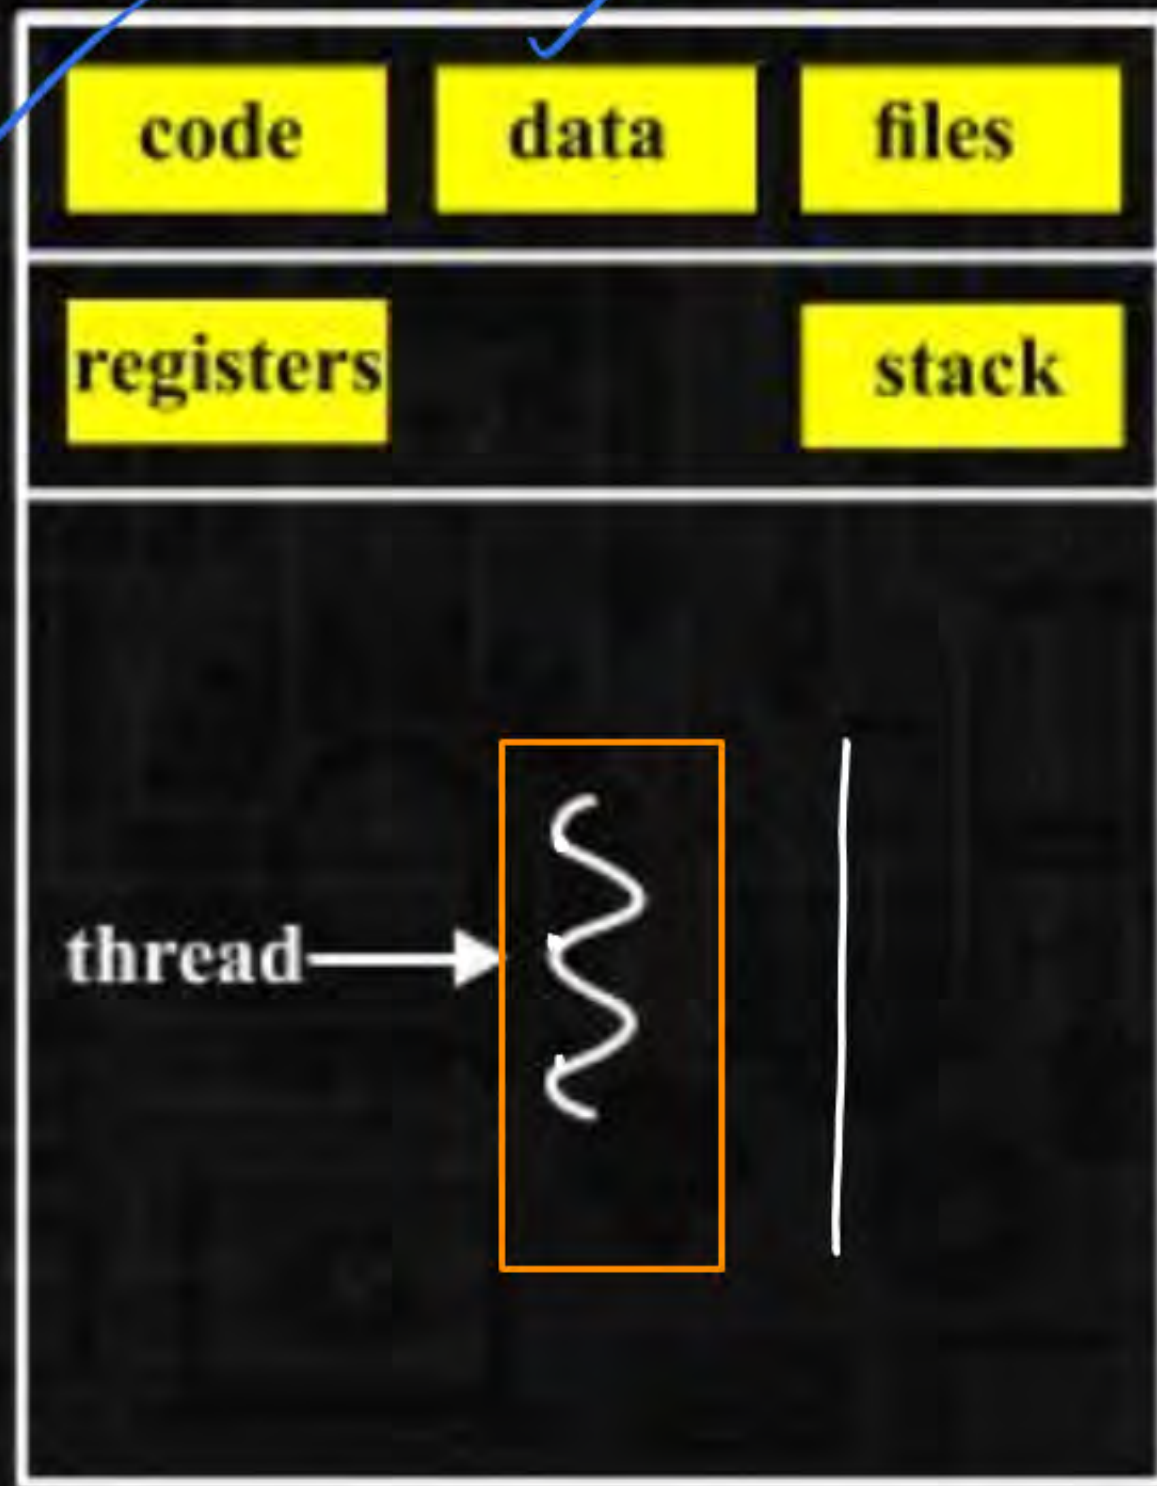
(Faster Context-switching)

→ is an Active Entity

→ is a Schedulable unit

→ is a unit of CPU utilization







single-threaded process

Traditional Process

Heavy wt. Process




multithreaded process



One Address Space


Memory





# Motivation :

## Client-Server Environment






Benefits :

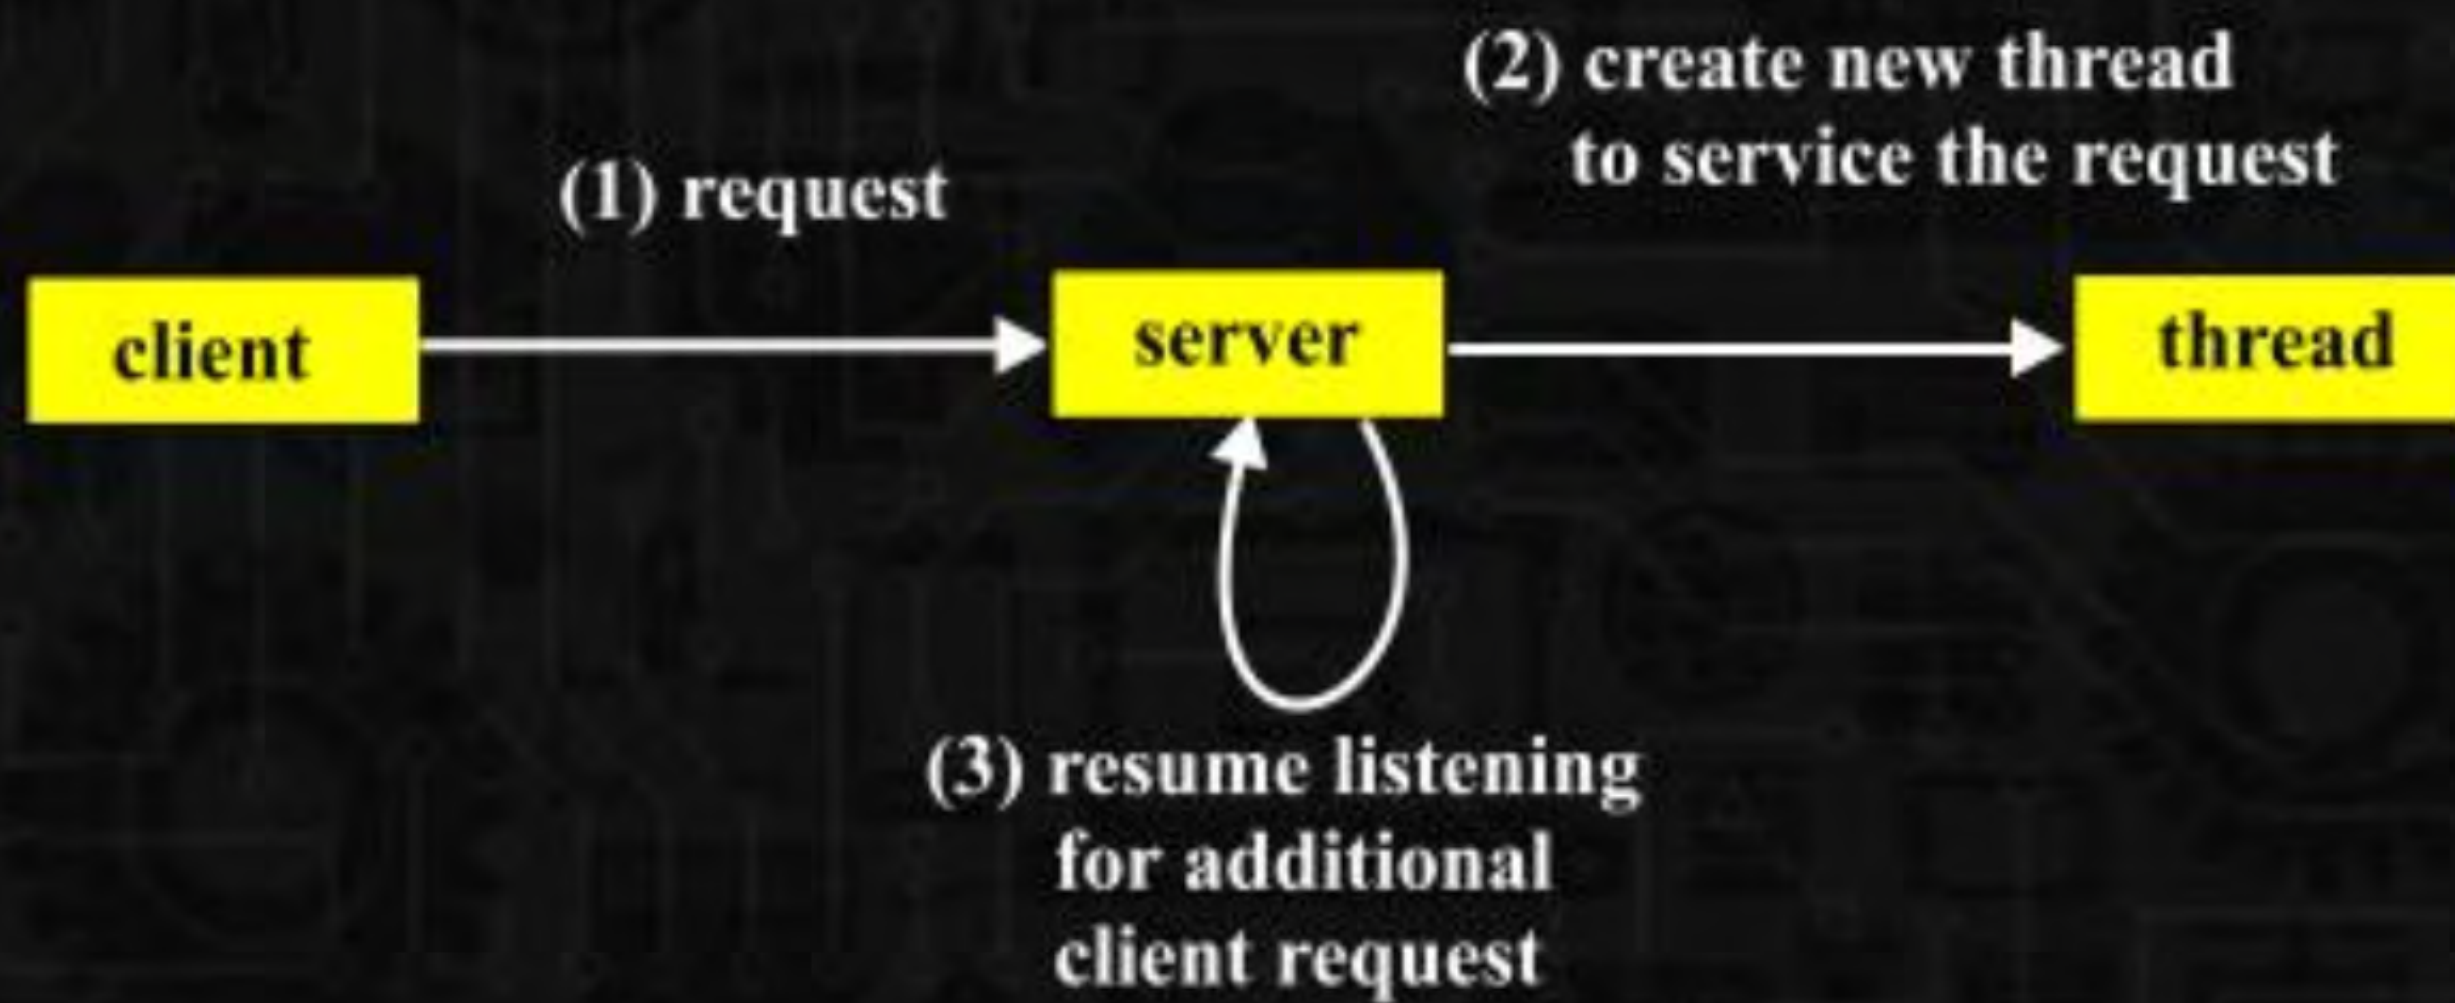
- 1) Resource Sharing  
(Address Space)
- 2) Economical design
- 3)  $\text{Size}_g(\text{TCB}) < \text{Size}_g(\text{PCB})$
- 4) Less Context-switch overhead  
(Thread-switching is faster)

5) Can exploit  
Multi-Core Architecture



6) Make computation faster







## The benefits of multithreaded programming can be broken down into four major categories:

1. **Responsiveness:** Multithreading an interactive application may allow a program to continue running even if part of it is blocked or is performing a lengthy operation, thereby increasing responsiveness to the user. This quality is especially useful in designing user interfaces. For instance, consider what happens when a user clicks a button that results in the performance of a time-consuming operation. A single-threaded application would be unresponsive to the user until the operation had completed. In contrast, if the time-consuming operation is performed in a separate thread, the application remains responsive to the user.



2. Resource sharing: Processes can only share resources through techniques such as. shared memory. and. message) passing.. .Such. techniques, must be explicitly arranged by the programmer. However, threads share the memory and the resources of the process to which they belong by default. The benefit of sharing code and data is that it allows an application to have several different threads of activity within the same address space.
3. Economy: Allocating memory and resources for process creation is costly. Because threads share the resources of the process to which they belong, it is more economical to create and context-switch threads. Empirically gauging the difference in overhead can be difficult, but in general it is significantly more time consuming to create and manage processes than threads. In Solaris, for example, creating a process is about thirty times




4. Scalability: The benefits of multithreading can be even greater in a multiprocessor architecture, where threads may be running in parallel on different processing cores. A single-threaded process can run on only one processor, regardless how many are available. We explore this issue further in the following section.



# Single Threaded Process



- ❑ So, far we have studied single threaded programs
- ❑ Recap: process execution
  - ❖ Pc points to current instruction being run
  - ❖ SP points to stack frame of current function call
- ❑ A program can also have multiple threads of execution
- ❑ What is a thread

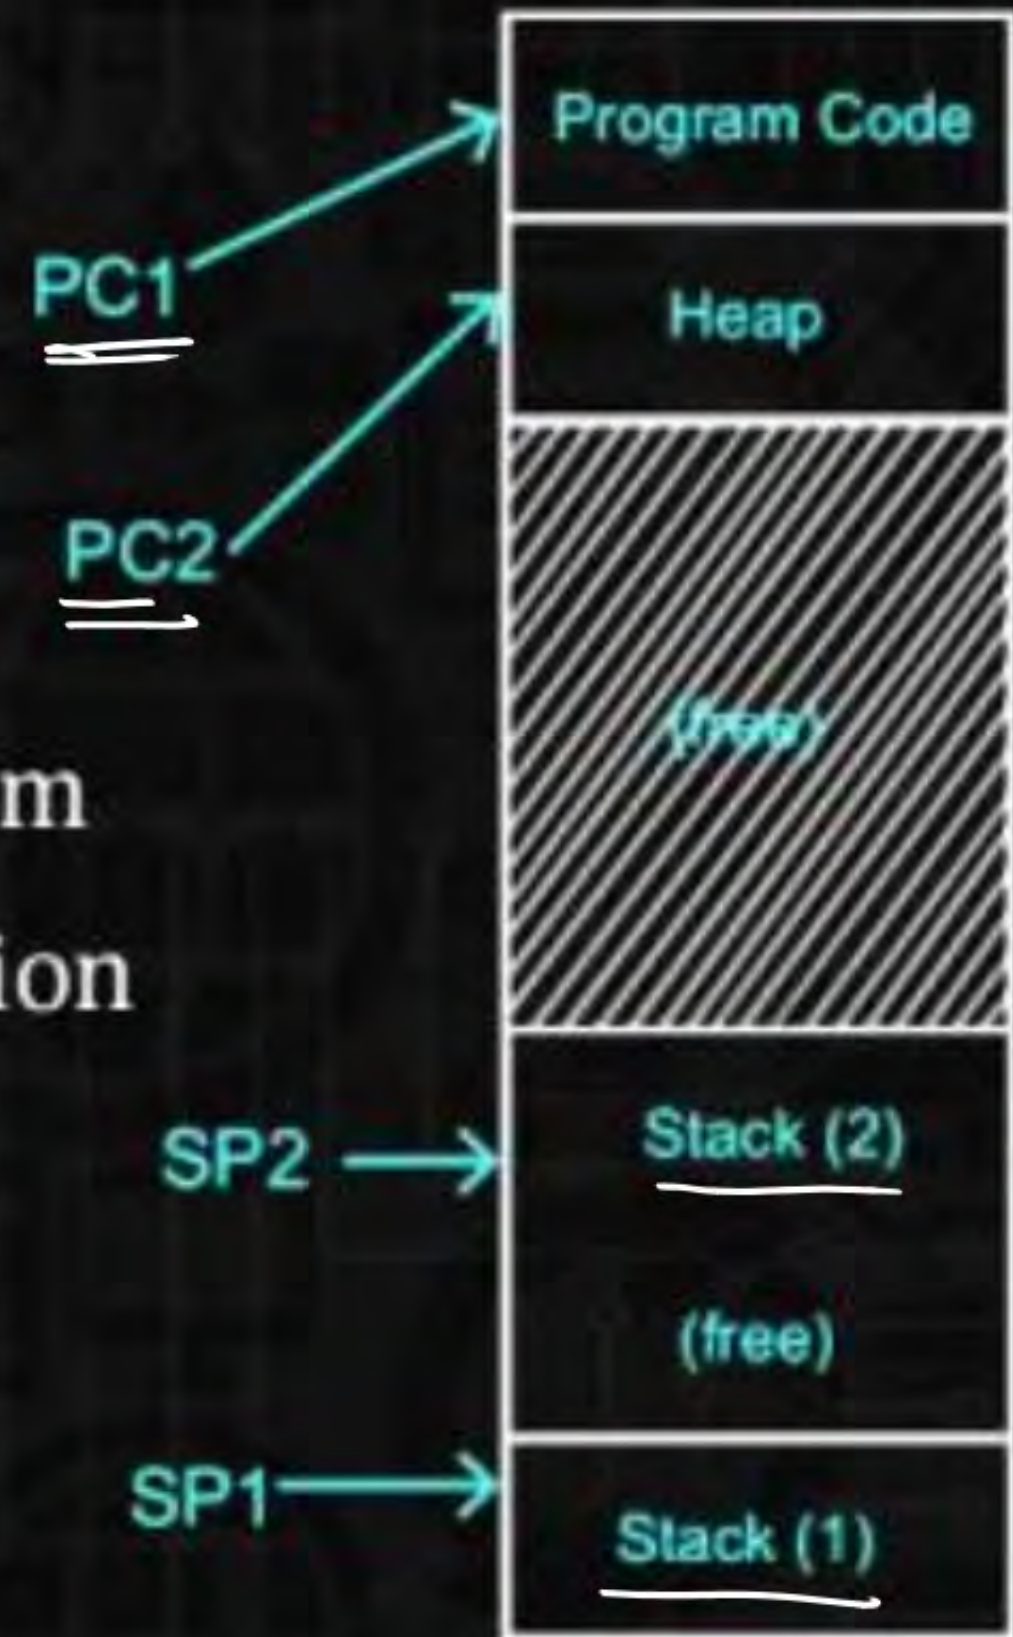




# Multi Threaded Process



- ❑ A thread is like another copy of a process that executes Independently.
- ❑ Threads share the same address space ( code heap)
- ❑ Each thread has a separate PC
  - ❖ Each thread may run over different part of a program
- ❑ Each thread has a separate stack for independent function calls





## Process Vs. Thread




- ❑ Parent P forks a child C
  - ❖ P and C does not share any memory
  - ❖ Need complicated IPC mechanism to communicate
  - ❖ Extra copies of code, data in memory
- ❑ Parent p execute two threads T1 and T2
  - ❖ T1 and T2 share parts of the address space
  - ❖ Global Variables can be used for communication
  - ❖ Smaller memory footprint
- ❑ Threads are like separate processes, except they share the same address space



# Threads



- ❑ Separate stream of execution within a single process
- ❑ Threads in a process not isolated from each other
- ❑ Each thread States (thread control block) contains
  - ❖ Registers including (EIP, ESP)
  - ❖ Stack





# Threads Vs Processes




- ❑ A thread has no data segment or heap.
- ❑ A thread cannot live on its own it need to be attached to a process.
- ❑ There can be more than one thread in a process. Each thread has its own stack.
- ❑ If a thread dies, its stack is reclaimed.
- ❑ A process has code, heap, stack, and other segments
- ❑ A process has at-least one thread.
- ❑ ~~Heads~~ <sup>Threads</sup> within a process share the same code, files.
- ❑ If a process dies, all thread die.



# Types of Threads:

User-level  
Threads  
(ULTs)

Kernel-  
level  
Threads  
(KLTs)



Multi-Threaded  
Kernel

(PThread)



## U.L.T's :

→ Threads that are created & Managed @ user-level without any support of O.S (Knowledge)

→ Flexibility

→ Transparency

→ Fastest Context Switching

(No Need of Mode Shifting)

→ Requirement of IO/Sys-Call exec, in a user-level Thread would result in blocking of the whole process,



- ❑ Create a thread in a process

Int

```
pthread_create(pthread_t*thread,  
Const pthread_attr_t*attr,  
void*(*start_routine) (void*),  
void*arg);
```

- ❑ Destroying a thread

void

```
pthread_exit(void*retval);
```

Thread identifier (TID) much like  
Pointer to a function which starts  
execution in a different thread

Arguments to the function

Exit value of the thread



## Pthread library contd.

- ❑ Join : wait for a specific thread to complete  
Int pthread\_join(pthread\_t thread, void\*\*retval);  
TID of the thread to wait for exit status of the thread



## Who manages threads?



Two strategies

- ❑ User threads


- ❖ Thread management done by user level threads library. Kernel knows nothing about the threads.

- ❑ Kernel threads

- ❖ threads directly supported by Kernel.
- ❖ Known as light weight processes.



## Use Level threads





### ❑ Advantages:

Scheduler can decide to give more time to a process having small number of threads.


Kernel-level threads are especially good for applications that frequently block.

### ❑ Disadvantages:

The kernel-level are slow (they involve kernel invocations .

Overheads in the kernel. (Since kernel must manage and schedule threads as well as processes. It required a full thread control block (TCB) for each thread to maintain information about threads.)







**Q. 1**

Which of the following is/are shared by all the threads in a process?

- I. Program counter
- II. Stack
- III. Address space
- IV. Registers

- A. I and II only
- B. III only ✓
- C. IV only
- D. III and IV only



Q. 2

Threads of a process share

- A. Global variables but not heap
- B. Heap but not global variables
- C. Neither global variables nor heap
- D. Both heap and global variables ✓



Q. 3

Which one of the following is FALSE?

- A. User level threads are not scheduled by the kernel. T
- B. When a user level thread is blocked, all other threads of its process are blocked. T
- C. Context switching between user level threads is faster than context switching between kernel level threads. T
- D. Kernel level threads cannot share the code segment. F



**Q. 4**

A thread is usually defined as a light weight process because an Operating System (OS) maintains smaller data structure for a thread than for a process. In relation to this, which of the following statement is ~~correct?~~ FALSE?

msQ

- A. OS maintains only scheduling and accounting information for each thread. ✗
- B. OS maintains only CPU registers for each thread. ✗
- ☒ C. OS does not maintain virtual memory state for each thread.
- ☒ D. OS does not maintain a separate stack for each thread. ✗



Q. 5

Consider the following statements about user level threads and kernel level threads. Which one of the following statements is FALSE?

- A. Context switch time is longer for kernel level threads than for user level threads. T
- B. User level threads do not need any hardware support. T
- C. Related kernel level threads can be scheduled on different processor in a multi-processor system. T
- ~~D. Blocking one kernel level thread blocks all related threads. F~~



Q. 6

Which one of the following is NOT shared by the threads of the same process?

☒ A.

Stack

☐ B.

Address Space

☐ C.

File Descriptor Table (Resource)

☐ D.

Message Queue (Resource)

} I.P.C Mechanisms



Q. 7

Consider the following statements with respect to user-level threads and kernel-supported threads

- I. ☒ Context switch is faster with kernel-supported threads.
- II. ☒ For user-level threads, a system call can block the entire process.
- III. ☒ Kernel supported threads can be scheduled independently.
- IV. ☒ User level threads are transparent to the kernel.

Which of the above statements are true?

☒ A. II, III and IV only

☐ B. II and III only

☐ C. I and III only

☐ D. I and II only

FCFS



My official Telegram channel

<https://t.me/KhalaelSirPw>



